HDF5 & Blosc2

Unleashing the full potential of Blosc2 from HDF5

Francesc Alted / @FrancescAlted

The Blosc Development Team / @Blosc2

CEO [[ ]] ironArray / @ironArray

# The Blosc Development Team

Aleix Alcacer
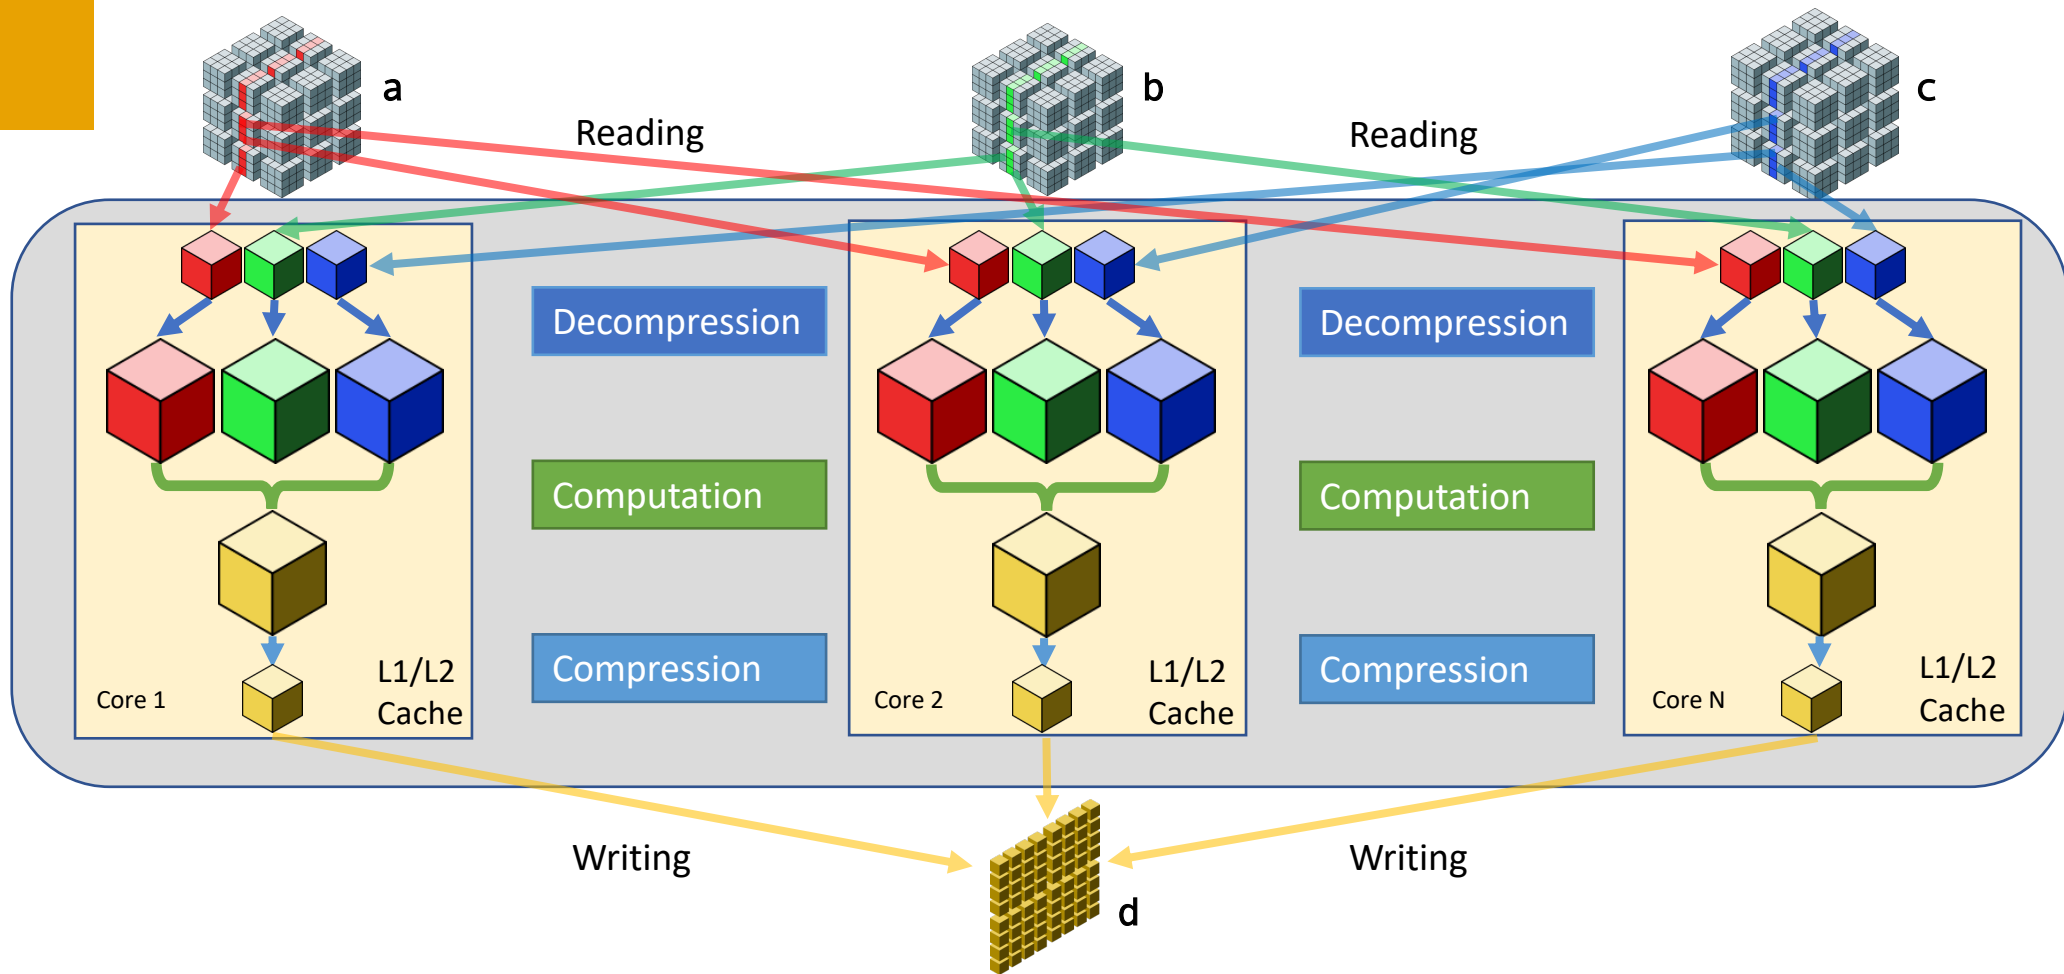
Oscar Guiñón

Marta Iborra

Alberto Sabater

Nathan Moinvaziri

Francesc Alted (BDFL)

A few words about ironArray

Intertwining compression and computation for improving performance

What is Blosc?

- ✓ Sending data from CPU to memory (and back) faster than *memcpy().*

- ✓ Split in blocks for better cache use: divide and conquer.

- ✓ It can use different filters (e.g. shuffle, bitsuffle) and codecs (e.g. LZ4, Zlib, Zstd, BloscLZ).

Binary dataset (Chunk)

L1/L2 Cache

CPU

Block 1
Block 2
...
Block N

Blosc Container
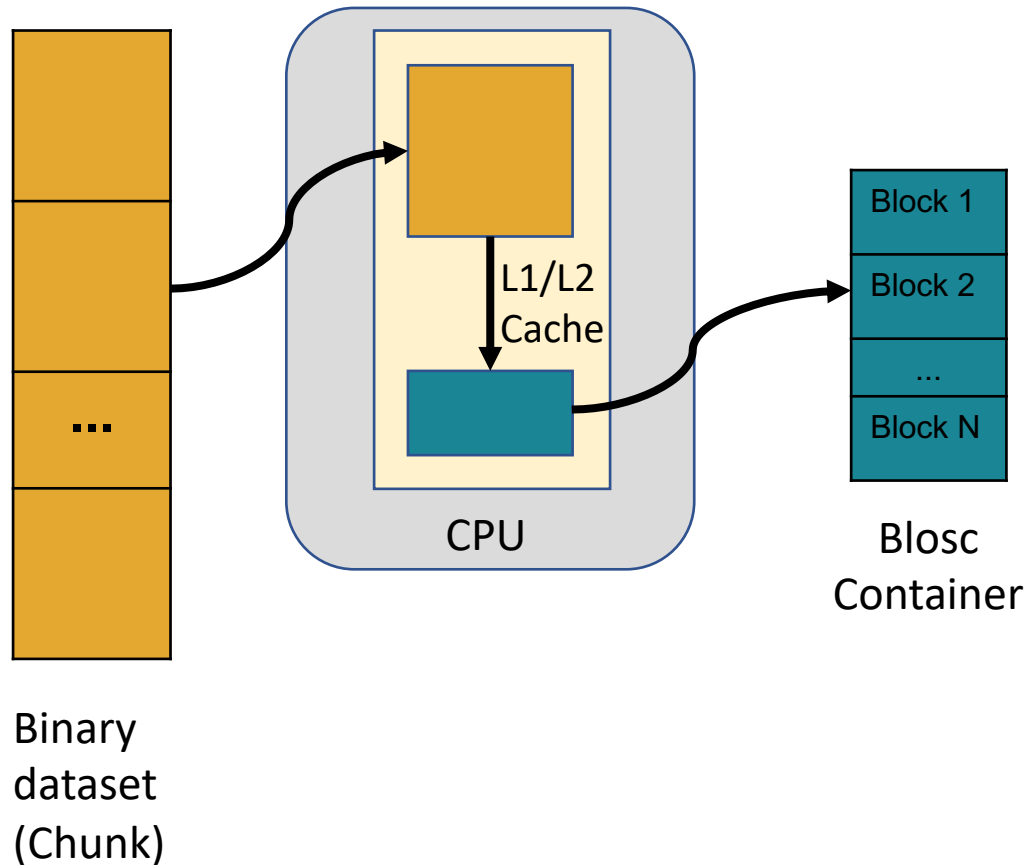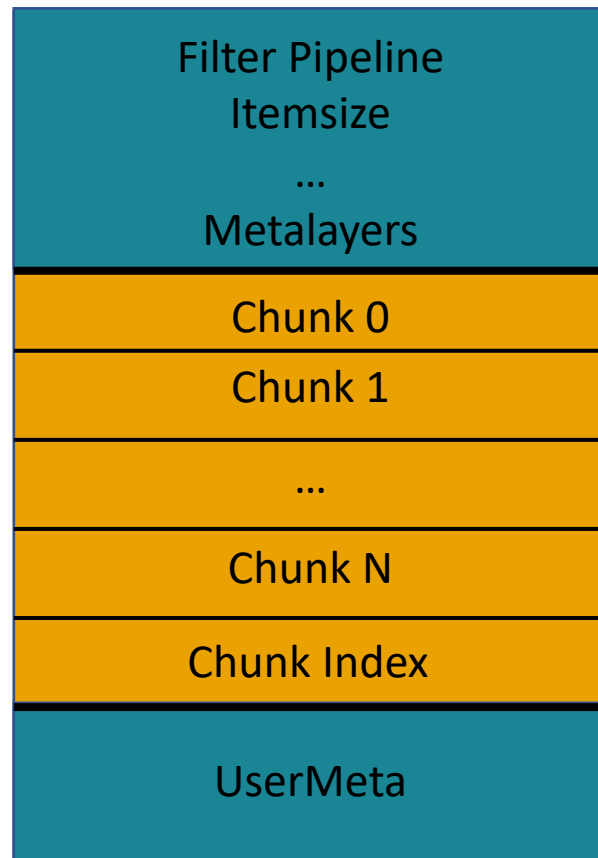
Origins of Blosc

- **2009**: it was very clear that compression was slowing down storage in PyTables/HDF5 a lot. Work began.

- **2010**: Blosc 1.0 was ready for production. Innovations:

  - Shuffle filter was optimized for SSE2 (*much* faster)
  - Multithreaded operation

- **2013**: Blosc gained multi-codec (LZ4, Snappy and Zlib where included)

- **2015**: hdf5-blosc plugin for HDF5 was released (hdf5plugin took over!)

- **2021**: Blosc2 appeared with **lots** of new features.

What is Blosc2?

- ✓ **Next generation of Blosc1**.

- ✓ New **63-bit frames** that expand over the existing 31-bit chunks in Blosc1.

- ✓ **Metalayers** for adding info for applications.

- ✓ Area for adding **metadata for users** (variable length).

Blosc2 Frame

Filter Pipeline Itemsize ... Metalayers
Chunk 0
Chunk 1
...
Chunk N
Chunk Index
UserMeta

Example of Decompression Speed

Summing up precipitation data (381.5 MB, float32)

https://www.blosc.org/posts/breaking-memory-walls/

HDF5: Multidimensions and Chunking

- Data can be stored in hypercubes, making retrieval very convenient.

- But there is a price to pay for this flexibility: HDF5 is known to be **slow** when writing/retrieving (hyperslabs of) data.

Direct Chunk Write/Read Feature

- Allow the aplication to handle the chunk I/O and bypass the powerful (but slow!) chunk handling machinery in HDF5.

- The result is that data can be handled up to about 10x faster (with efficient pre and post processing in the app).
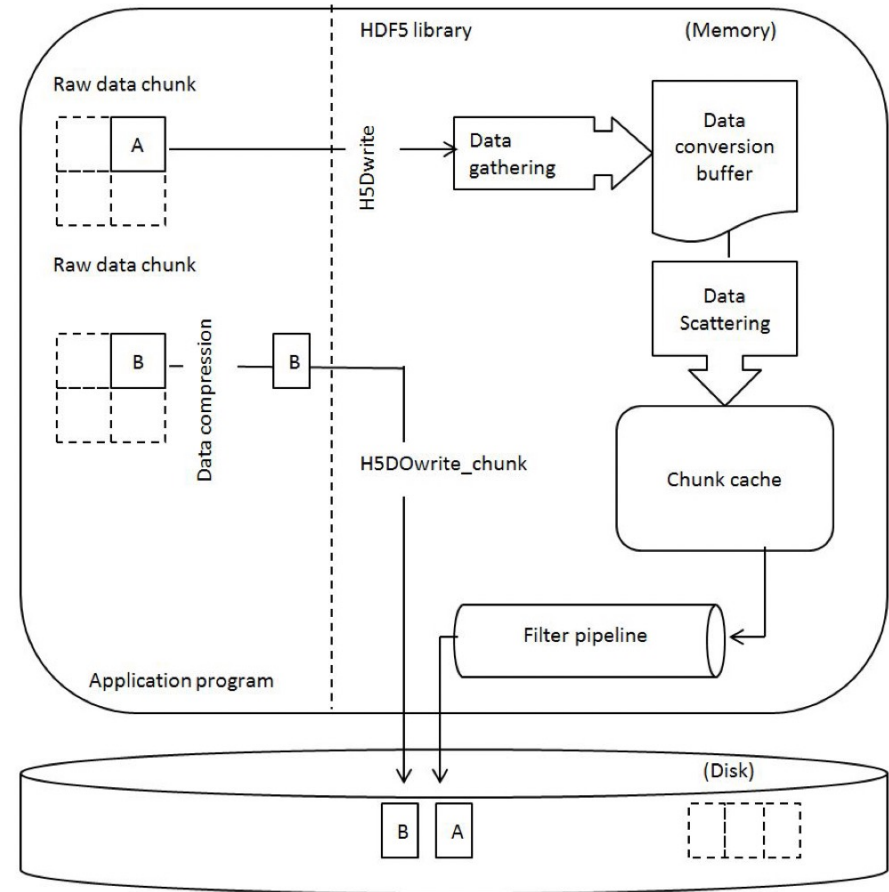
Proposal 1: Use Blosc2 Inside Direct Chunk

All compression and decompression executed in parallel via Blosc2!
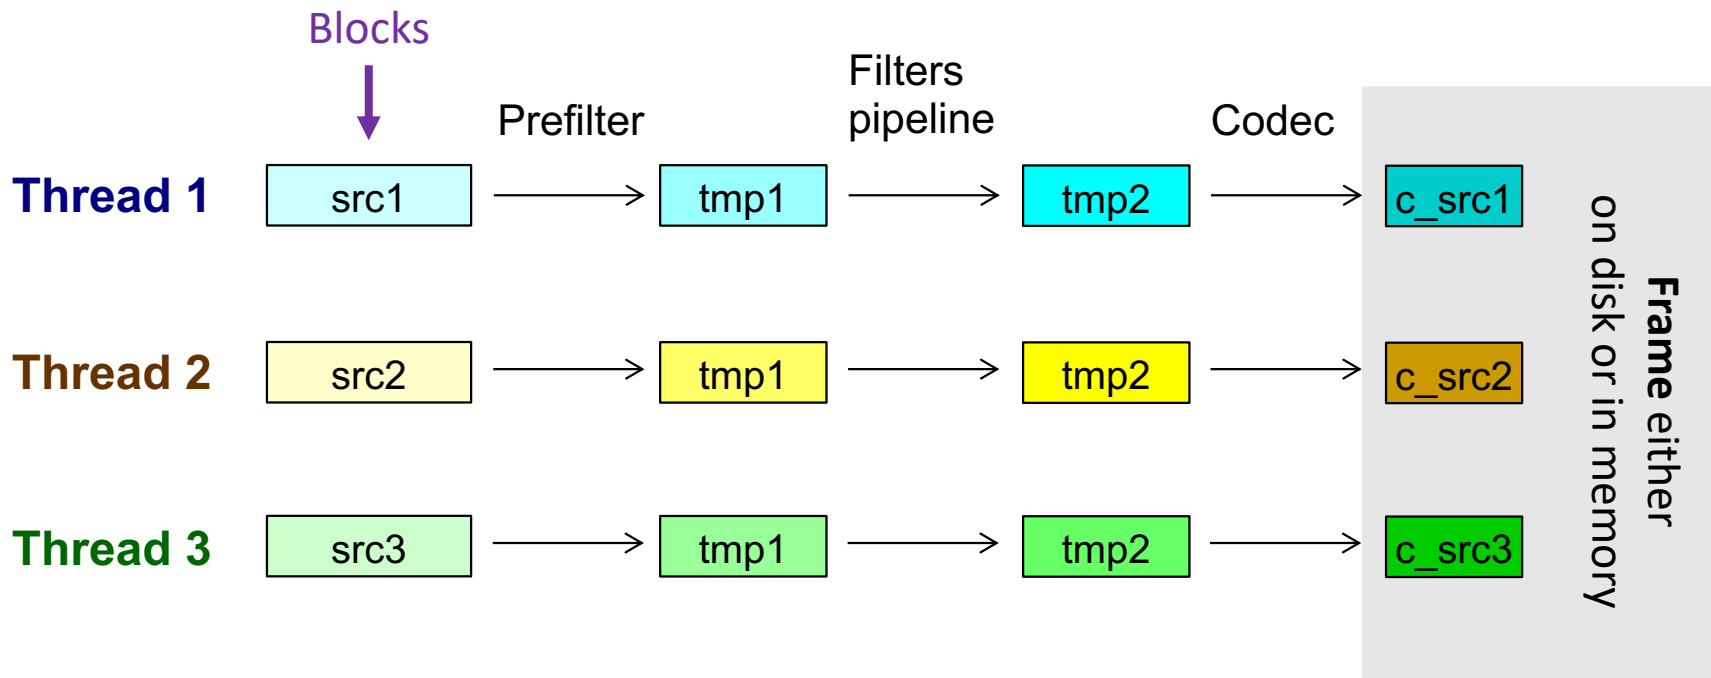

- HDF5 stores compressed data

- Blosc2 can do parallel I/O
- Blosc2 can do chunk reads with enhanced selectivity from disk

Blosc2 Advantages
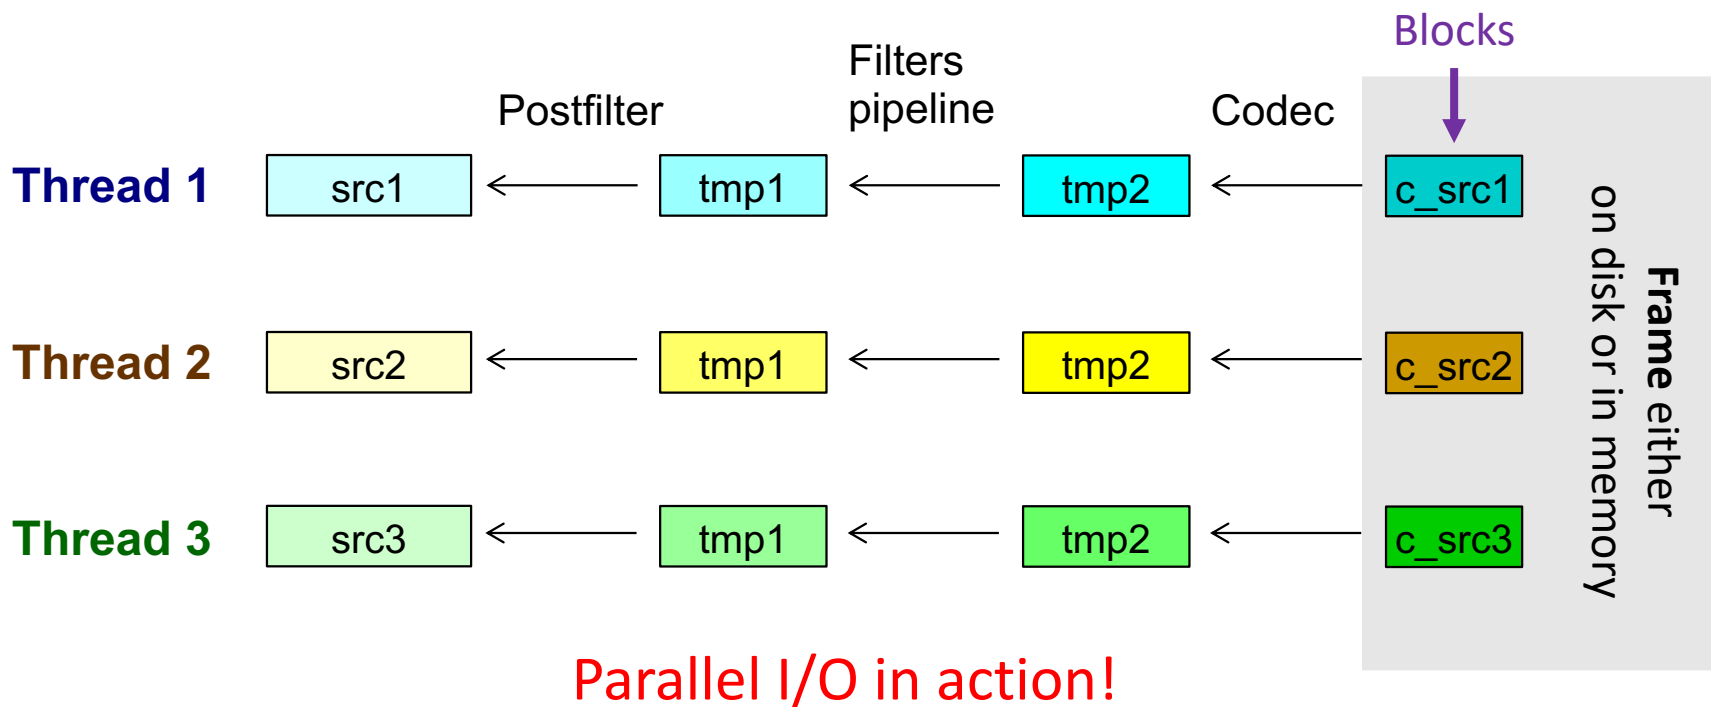
Blosc2: Fine Tuned Cache Usage

Compression: chunks are split in blocks for CPU cache sake


Buffers are reused **inside** CPU caches -> speed!

Blosc2: Leveraging I/O Parallelism

Decompression: blocks are read in parallel from storage


Parallel I/O in action!

Caterva: Blosc2 Goes Multidimensional

- ✓ Metalayer representing multidimensionality

- ✓ Each Caterva array is split in chunks

- ✓ Each chunk is split in blocks

- ✓ All the partitions are multidimensional!


Chunk

Array

Block

Benefits of the Caterva Layer

- Get **improved compression ratio** because data is packed in a way that can show higher spatial locality.

- Also, get **improved hyperslab query speed**, i.e. some blocks can be masked out so as to not read them.

ZFP: a new registered plugin

# Masked & Paralel I/O in Multidim Datasets

Much more selective and faster queries!

Caterva (https://github.com/Blosc/caterva) and ironArray (https://ironarray.io)

Masked & Paralel I/O in Multidim Datasets

Slicing Performance on disk (with an optimized dimension)

Better performance in general
(except for dimensions where retrieving a chunk is already optimal)

https://ironarray.io/docs/html/tutorials/03.Slicing_Datasets_and_Creating_Views.html

Blosc2: Paralellism and Efficiency

Mean of 3 arrays of 3 GB each (on disk)

- In the plot: 3 compressed arrays are decompressed, operated with, and the result is compressed again.

- ironArray is using Blosc2.

- When handled correctly, parallelism can buy not only speed, but also less memory resources!
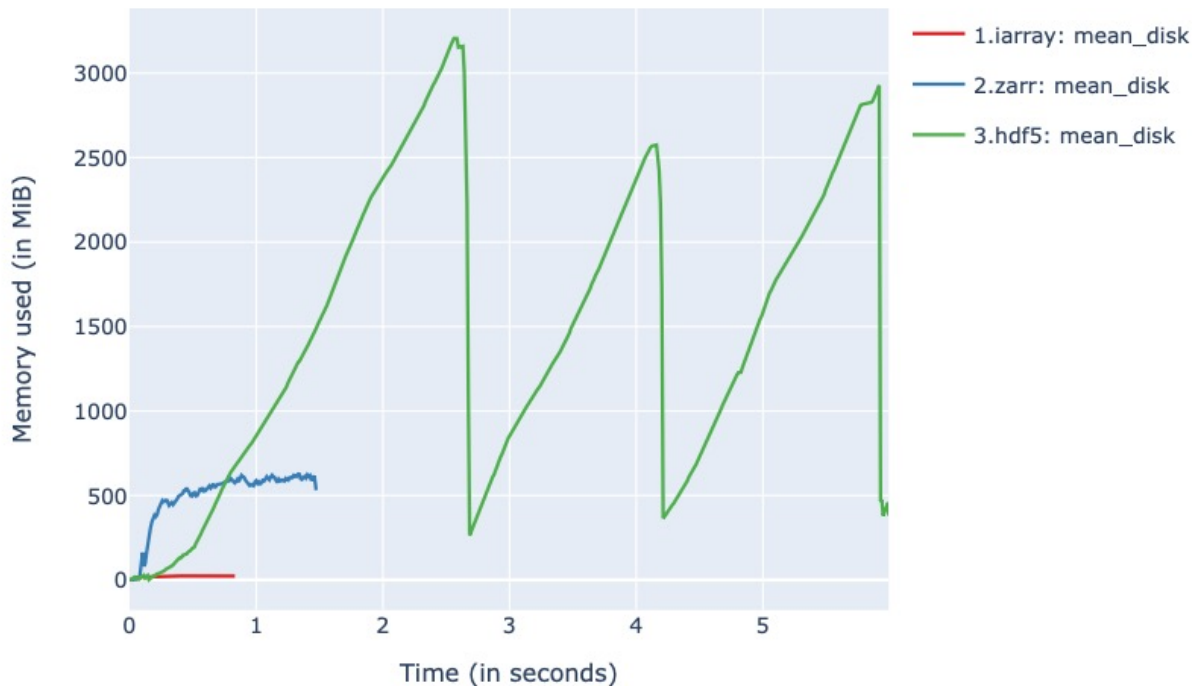
https://ironarray.io/docs/html/bench/05.Reductions_OnDisk.html

Blosc2 + HDF5 Direct Chunk Quick Benchmark

Quick test on data from ERA5 dataset, using different measurements (wind, temp, precip...). The datasets are ~3 GB (uncompressed) each.

- Blosc2 + HDF5 speed-up is typically between 10x and 20x for writing.

- Blosc2 + HDF5 speed-up is typically between 20x and 30x for reading.

Note: this is very preliminary, but the measuments make sense, specially when using a fast storage (in this case OS FS cache).

Blosc2 + H5Dchunk_write

Blosc2 + H5Dchunk_read

Adaptability: Plugins in Local Registry

Filters registry

BLOSC_SHUFFLE	1
BLOSC_BITSHUFFLE	2
BLOSC_DELTA	3
. . .	
BLOSC_NDCELL	32
BLOSC_NDMEAN	33
. . .	
urfilter1	160
urfilter2	161
. . .	

☐ Blosc official registered filters
☐ User local filters

User defined filter:

```
int urfilter2(
    blosc2_filter *filter) {
    …
}
```

To register locally:
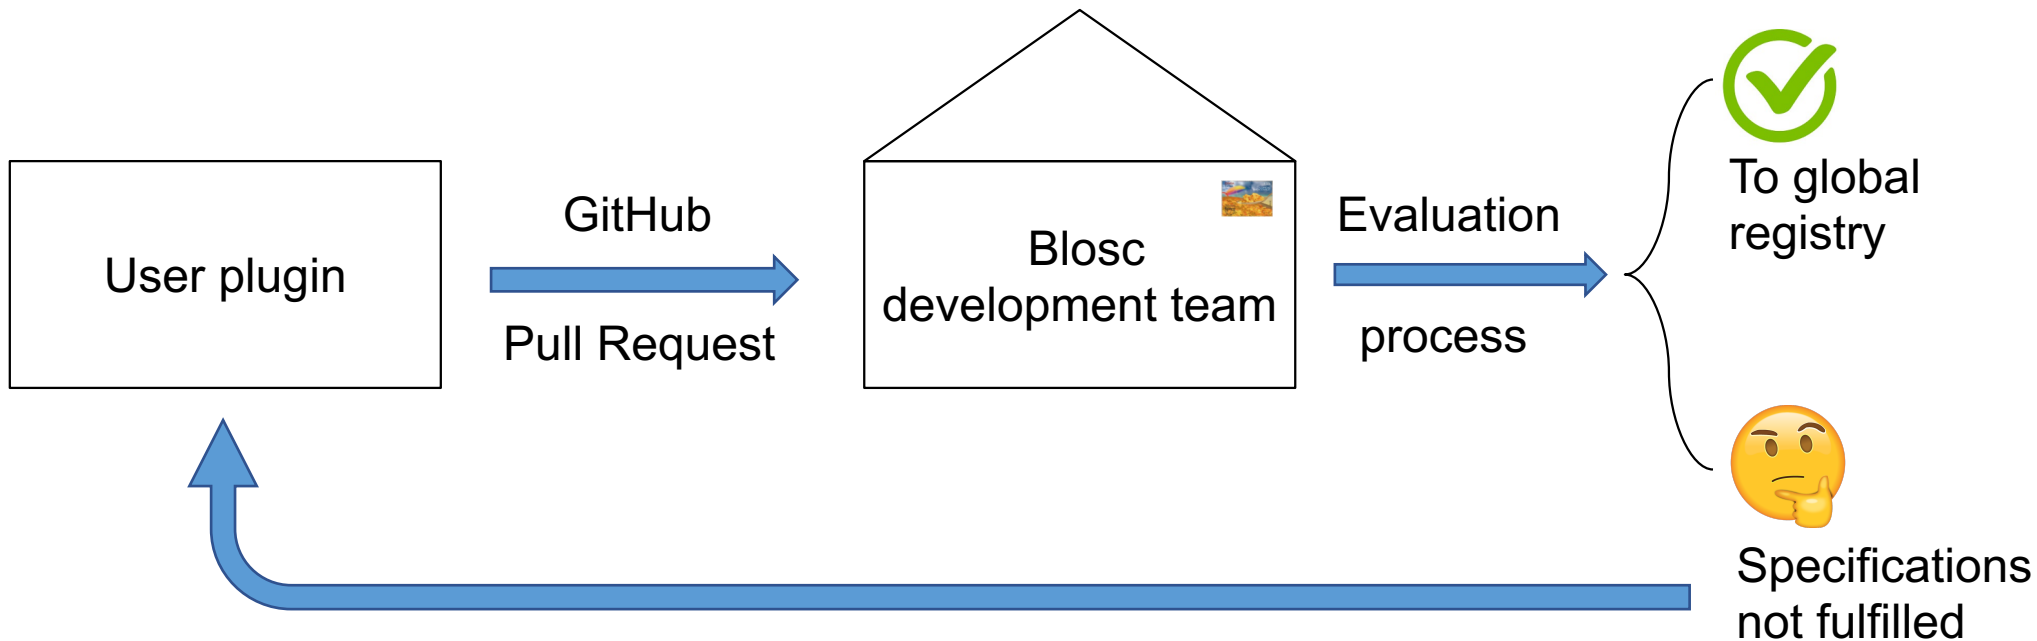
```
blosc2_register_filter(
    urfilter2)
```

Can be used now:

```
cparams.filters[4] = 161;
```

(Similar functionality to the plugin interface in HDF5)

Registering Global Plugins in Blosc2

User plugin

GitHub

Pull Request

Blosc development team

Evaluation process

To global registry

Specifications not fulfilled

Specs: https://github.com/Blosc/c-blosc2/blob/main/plugins/README.md

Proposal 2: Help in Determing Optimal Compression Pipelines

We are offering a service for adapting to the user data, and determining:

- Set of most useful codecs to be used

- Set of most useful filters to be used

We produce **specific versions** of **BTune**, a machine learning tool for selecting the best pipeline candidate on a **chunk by chunk** basis, that adapts to the needs of the user.

Fine Tuning Performance with BTune

- BTune can fine tune the different parameters of the underlying Blosc2 storage to perform as best as possible.

- Active during the compression pipeline. **Automatically learns the best parameters on the go**.


BTune State Diagram

Demo time

https://btune.blosc.org
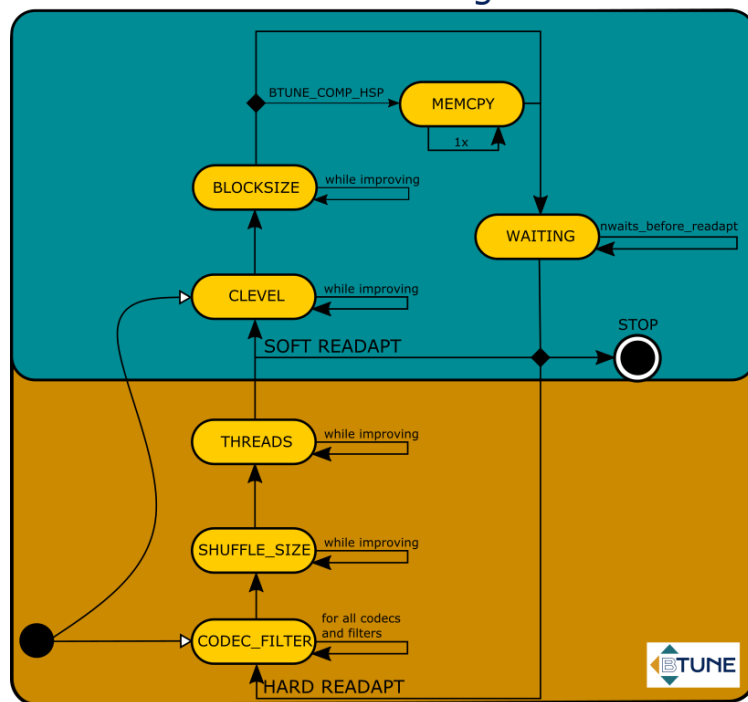
Work in Improving BTune

- Currently BTune needs some warm-up tests (hard and soft readapts) so as to come with a sensible guess.

- We are planning to shorten this warm-up period by using deep learning techiques.

- The idea is to come with some **predictor for the entropy** for every chunk and train a neural network. This will be used for reaching the **sensible guess faster**.

Current Investigation (Preliminary Results)

Codecs with similar features: Not good predictions

Codecs with different features: Much better predictions!

Conclusion

Blosc2 Helps Saving Resources

Blosc2 orchestrates a **rich set of codecs and filters** for:

- **CPU parallelization** via multithreading
- Reuse and sharing internal buffers for **optimal memory consumption**
- **Parallel I/O**
- **More selective** hyperslabs

The result is a highly efficient tool for
**compressing and accessing your data your way**

Summary of Proposals

1. Use Blosc2 in combination with HDF5 direct chunking mechanism for efficient compression and parallel I/O.

2. Help in determing optimal compression pipelines by adapting to user data and using machine learning techniques.

The Blosc team would be glad to be involved
in efforts towards these goals

Thanks to donors & contracts & contractors!

HUAWEI

ironArray

NUMFOCUS
OPEN CODE = BETTER SCIENCE

python SOFTWARE FOUNDATION

Google

ESRF
The European Synchrotron

Jeff Hammerbacher

Without them, we could not have possibly put Blosc2 into production status: Blosc2 2.0.0 came out in June 2021; now at 2.1.1.

Enjoy data!

https://blosc.org/