

Data Oriented Programming

or:

How I Learned to Stop Worrying about CPU Speed
and Love Memory Access

Francesc Alted
Freelance Trainer And Developer

PyData Conference, Berlin
July 26, 2014

Overview

- Motivation
- The Data Access Issue
 - Why Modern CPUs Are Starving
 - Why Caches?
 - Techniques For Fighting Data Starvation
- Optimal Containers for Big Data

About Me

- I am the creator of tools like PyTables, Blosc, bcolz, and a long-term maintainer of Numexpr
- I am an experienced developer and trainer in:
 - Python (almost 15 years of experience)
 - High Performance Computing and Storage
- Also available for consulting

Motivation

CPU Speed Is No Longer *The Holy Grail*

Suppose that we want to compute the next polynomial:

$$0.25x^3 + 0.75x^2 + 1.5x - 2$$

in the range $[-1, 1]$ with a step size of 2×10^{-7} in the x axis

...and want to do that as FAST as possible...

Using NumPy

```
import numpy as np  
  
N = 10*1000*1000  
  
x = np.linspace(-1, 1, N)  
  
y = .25*x**3 + .75*x**2 - 1.5*x - 2
```

That takes around 1.60 sec on some machine
(Intel Xeon E5520 @ 2.3 GHz). How to make it faster?

'Quick & Dirty' Approach: Parallelize

- Computing a polynomial is “embarrassingly” parallelizable: just divide the domain to compute in N chunks and evaluate the expression for each chunk.
- This can be easily implemented in Python by using the multiprocessing module (so as to bypass the GIL).
- Using 2 cores, the 1.60 sec is reduced down to 1.18 sec, which is a 1.35x improvement. Not bad.
- We are done! Or perhaps not?

A Better Approach: Optimize

The NumPy expression:

$$(I) \ 0.25x^3 + 0.75x^2 + 1.5x - 2$$

can be rewritten as:

$$(II) ((0.25x + 0.75)x + 1.5)x - 2$$

- Exec time goes from 1.60 sec to 0.30 sec
- Much faster (4x) than using two processors with the multiprocessing approach (1.18 sec).

First Take Away Message

- Do not blindly try to parallelize right away: Optimizing normally gives better results

And a serial codebase is normally much easier to code and debug!

Use numexpr

Numexpr is a JIT compiler, based on NumPy, that optimizes the evaluation of complex expressions. Usage is simple:

```
import numpy as np
import numexpr as ne
N = 10*1000*1000
x = np.linspace(-1, 1, N)
y = ne.evaluate(".25*x**3 + .75*x**2 - 1.5*x - 2")
```

This takes 0.14 sec to complete (11x faster than the original NumPy: 1.60 sec)

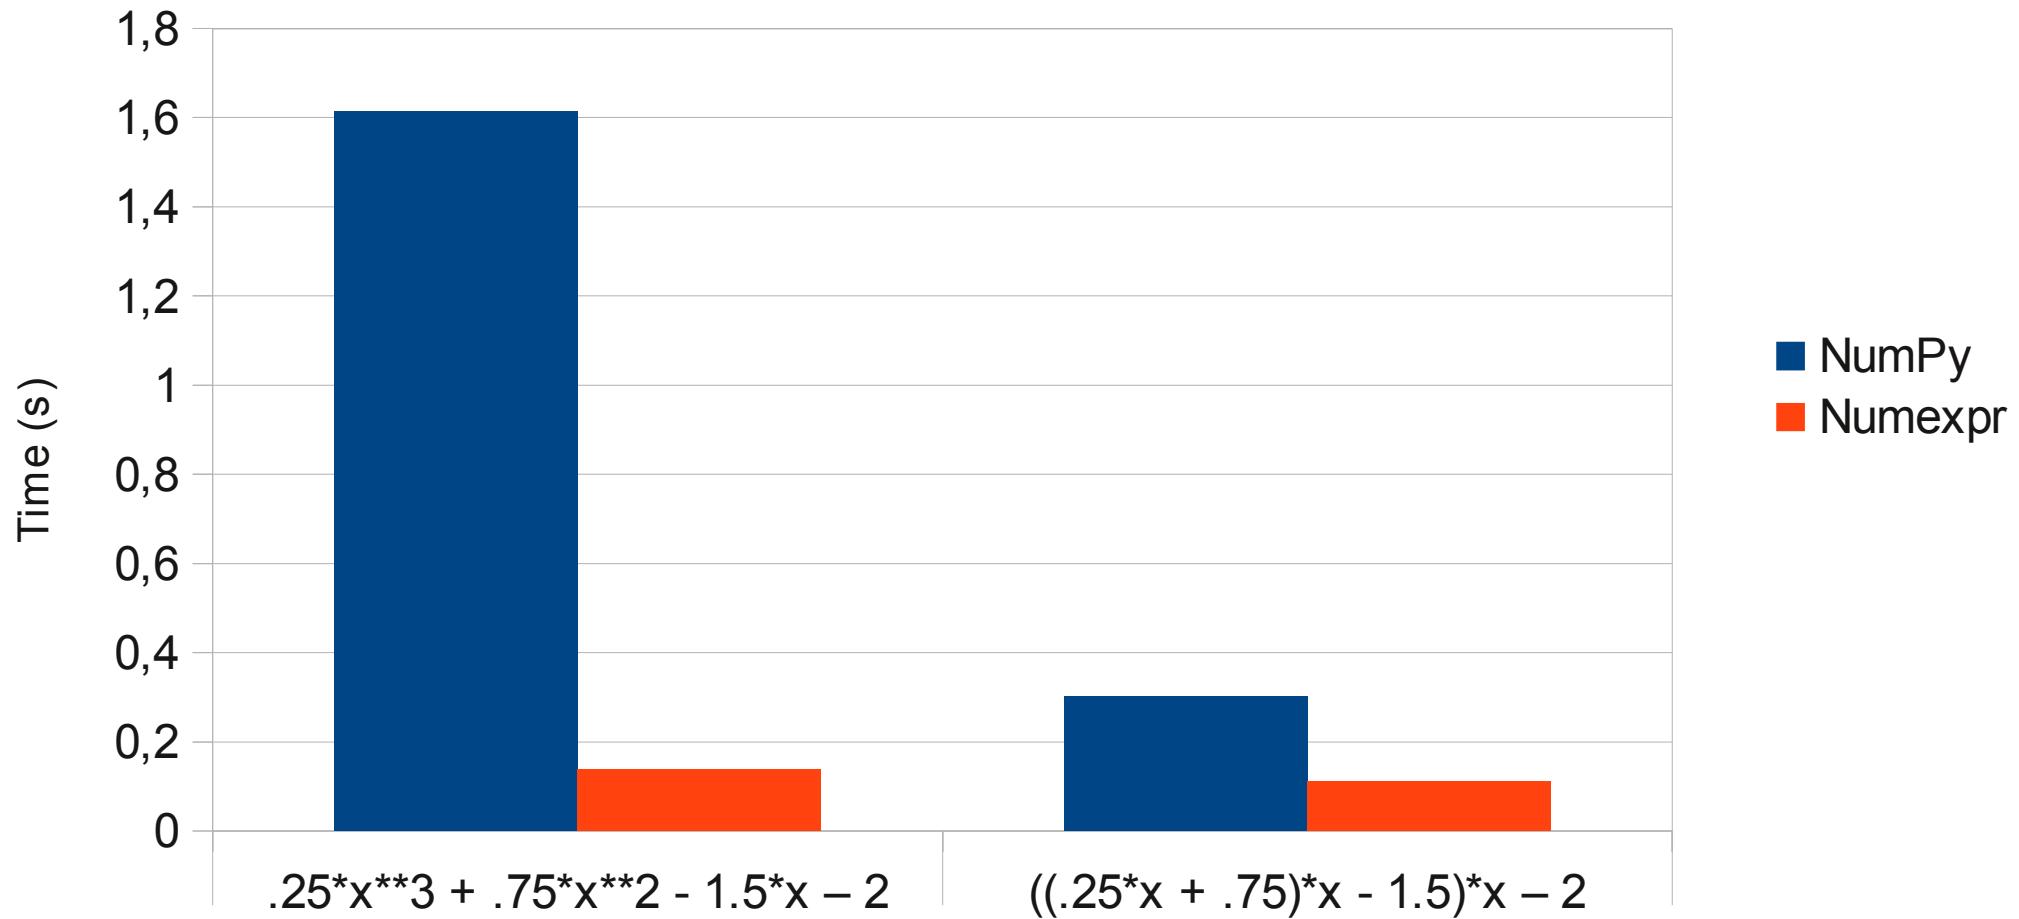
Fine Tuning numexpr

Numexpr is also sensible to computer-friendly expressions like:

$$(II) ((0.25x + 0.75)x + 1.5)x - 2$$

- This takes 0.11 sec (3x faster than NumPy)
- 0.14 sec were needed for the original expression, that's a 25% faster

Time to evaluate polynomial (1 thread)



Power Expansion

Numexpr expands expression:

`0.25*x3 + 0.75*x**2 + 1.5*x - 2`**

to:

`0.25*x*x*x + 0.75*x*x + 1.5*x - 2`

so, no need to use the expensive `pow()`

One Remaining Question

Why numexpr can execute this expression:

$$((0.25x + 0.75)x + 1.5)x - 2$$

3x faster, even using a single core?

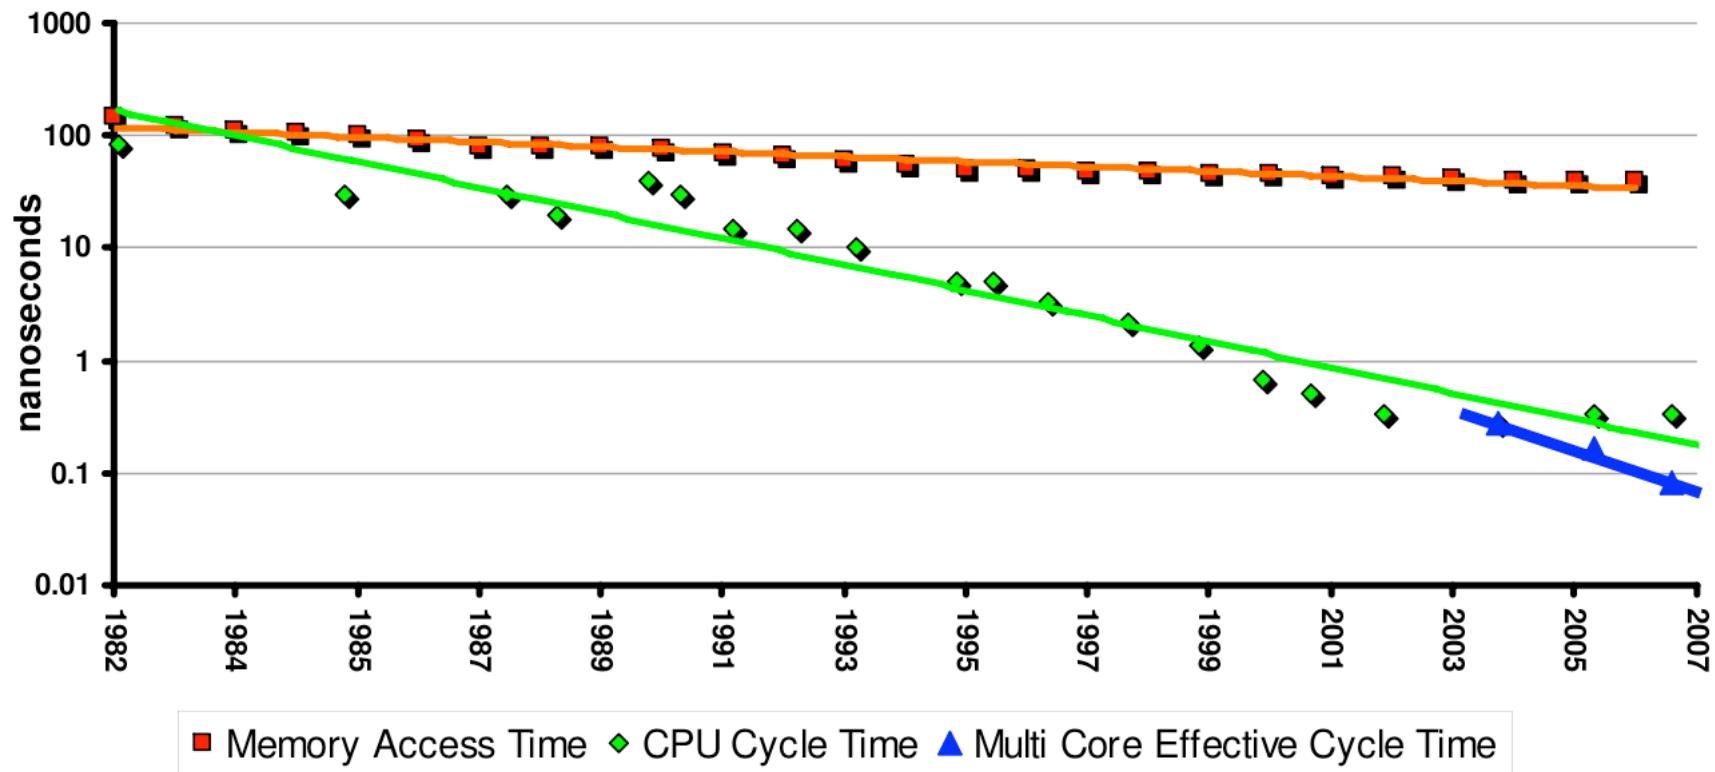
Short answer: making a more efficient use of the memory resource

The Starving CPU Problem

The Starving CPU Problem

- Current CPUs typically stay bored, doing nothing most of the time
- Why so?
- Because they are basically waiting for data

Memory Access Time vs CPU Cycle Time





The Memory System

*You Can't Avoid It,
You Can't Ignore It,
You Can't Fake It*

Book in 2009

Bruce Jacob

***SYNTHESIS LECTURES ON
COMPUTER ARCHITECTURE***

Mark D. Hill, *Series Editor*

Copyrighted Material

The Status of CPU Starvation in 2014

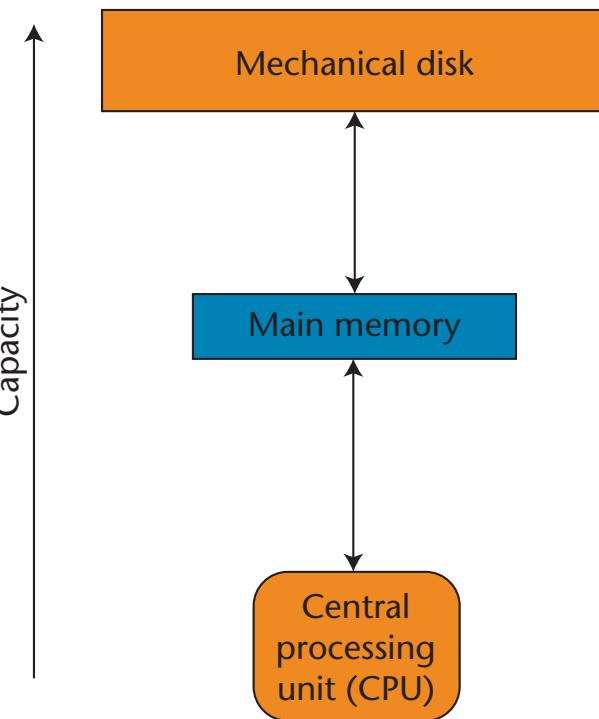
- Memory latency is much slower (between 100x and 500x) than processors.
- Memory bandwidth is improving at a better rate than memory latency, but it is also lagging behind processors (between 30x and 100x).

CPU Caches to the Rescue

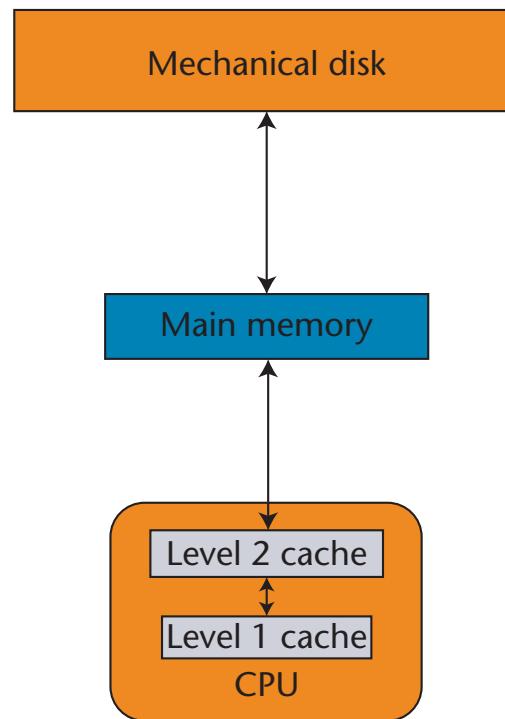
- CPU cache latency and throughput are much better than memory
- However: the faster they run the smaller they must be

CPU Cache Evolution

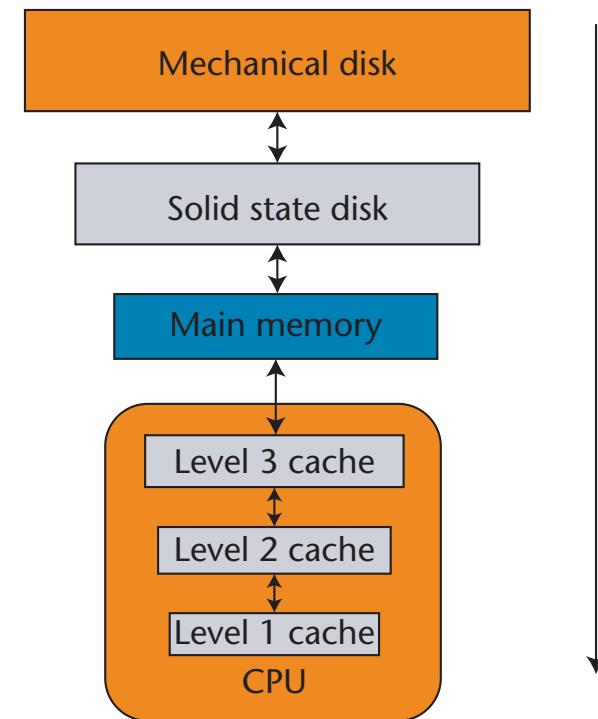
Up to end 80's



90's and 2000's



2010's



Capacity ↑

Speed ↓

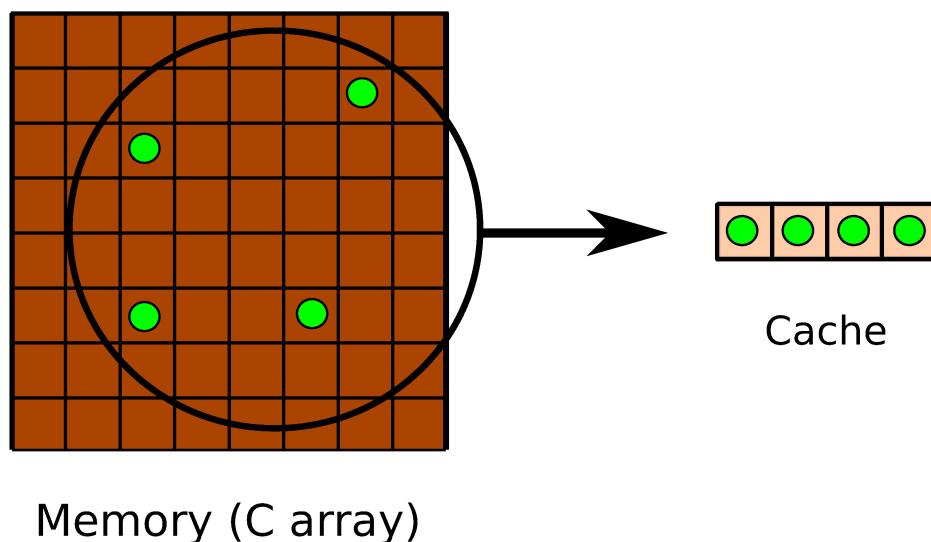
When CPU Caches Are Effective?

Mainly in a couple of scenarios:

- Time locality: when the dataset is reused
- Spatial locality: when the dataset is accessed sequentially

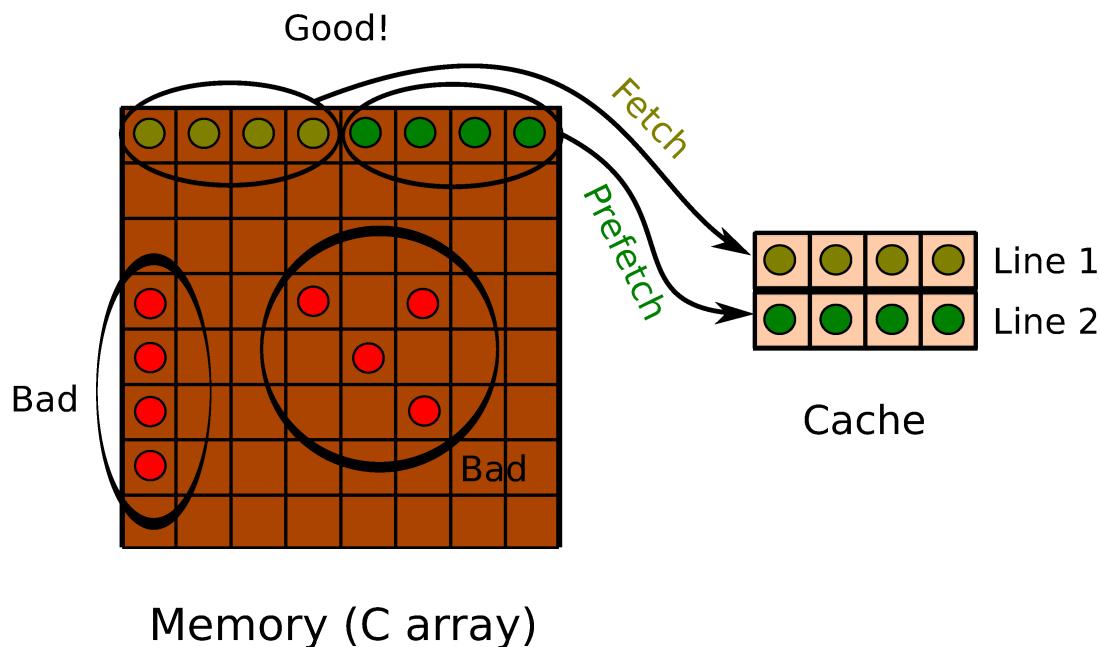
Time Locality

Parts of the dataset are reused



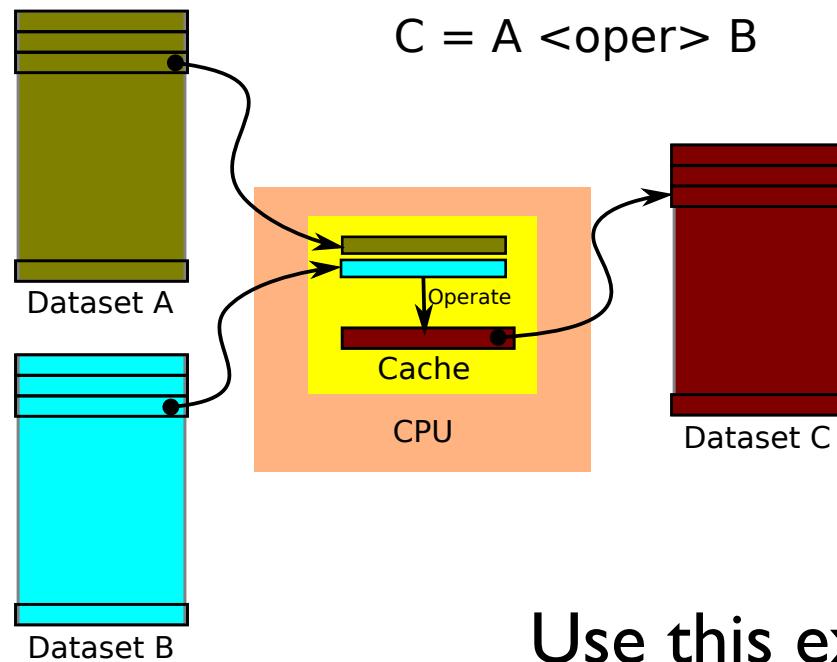
Spatial Locality

Dataset is accessed sequentially



The Blocking Technique

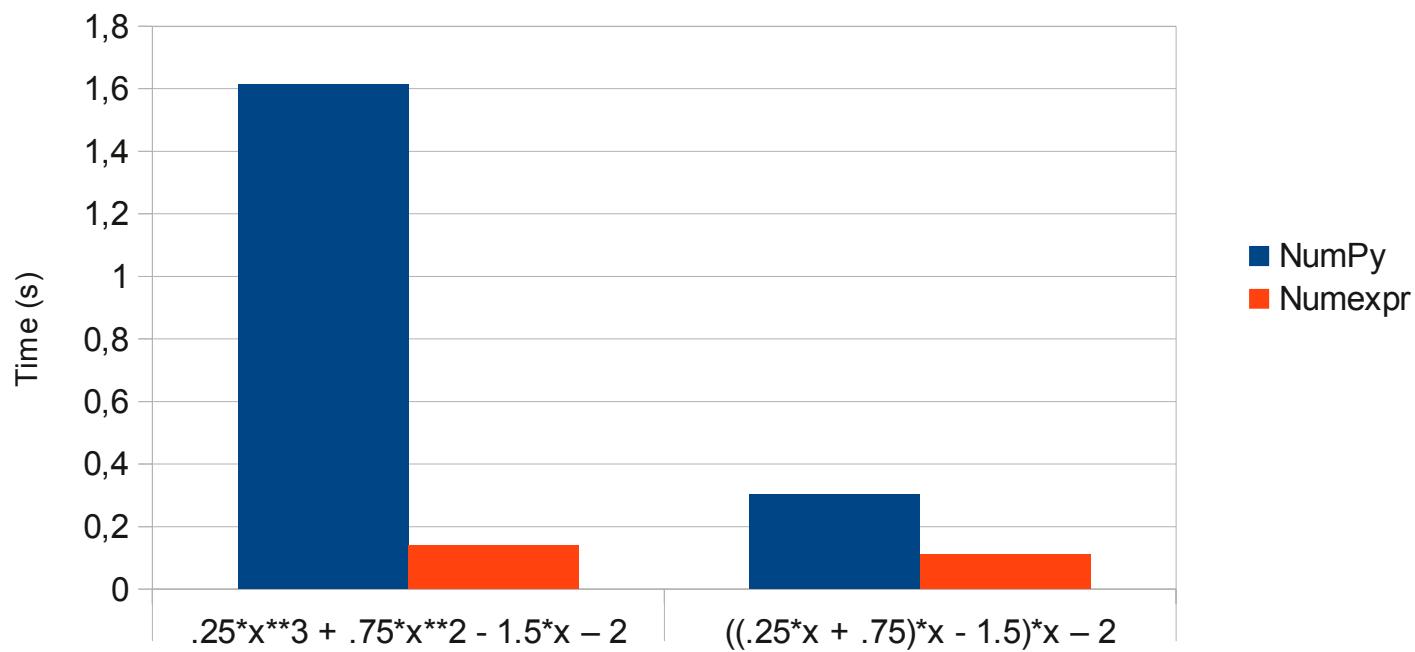
When accessing disk or memory, get a **contiguous** block that fits in CPU cache, operate upon it and **reuse** it as much as possible.



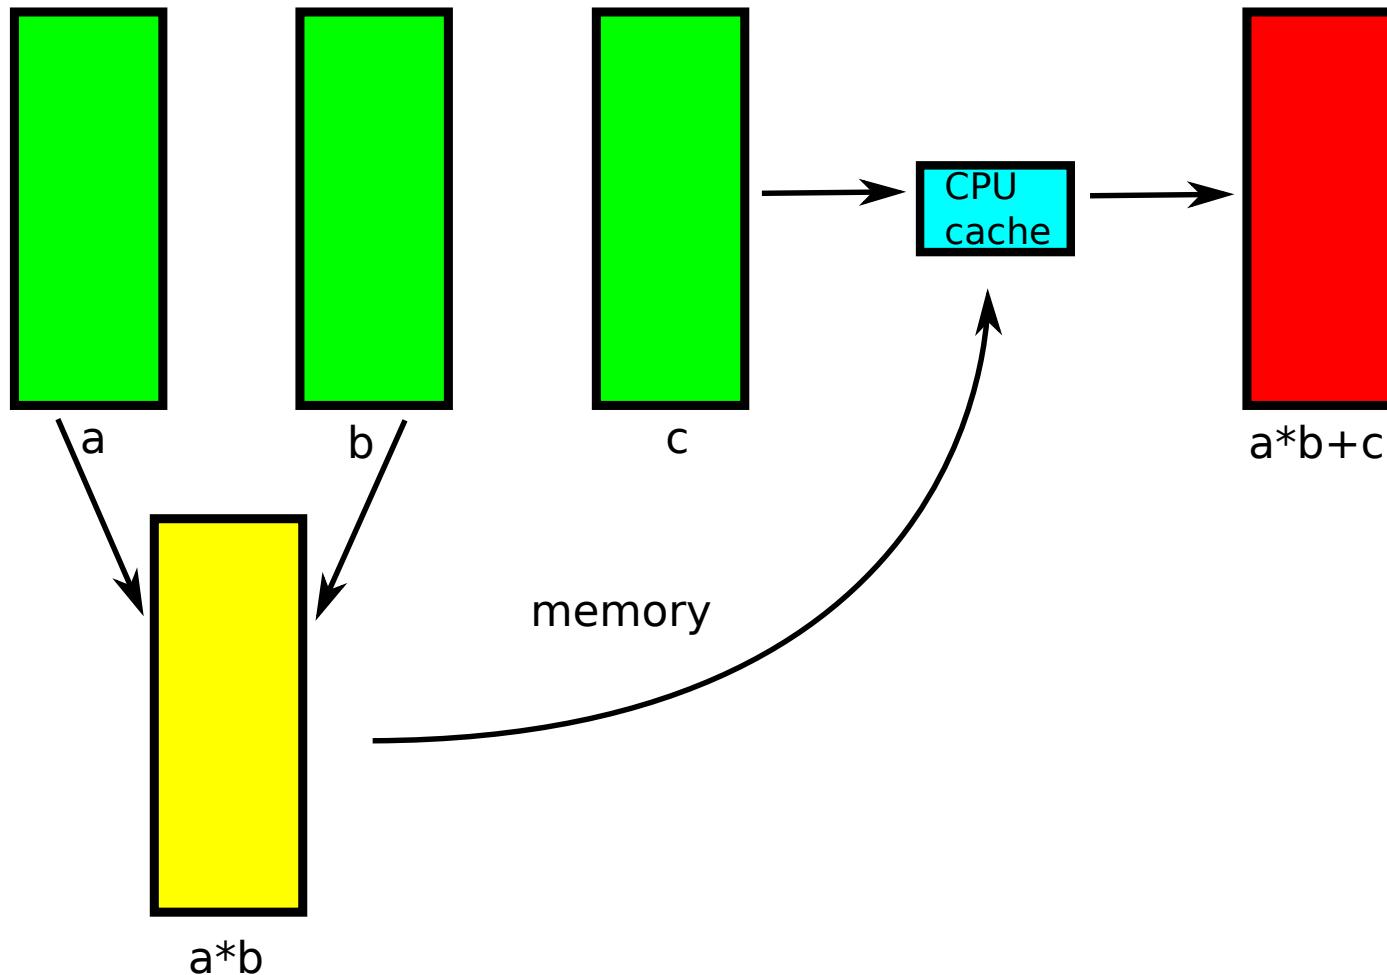
Use this extensively to leverage
spatial and **temporal** localities

Time To Answer Pending Questions

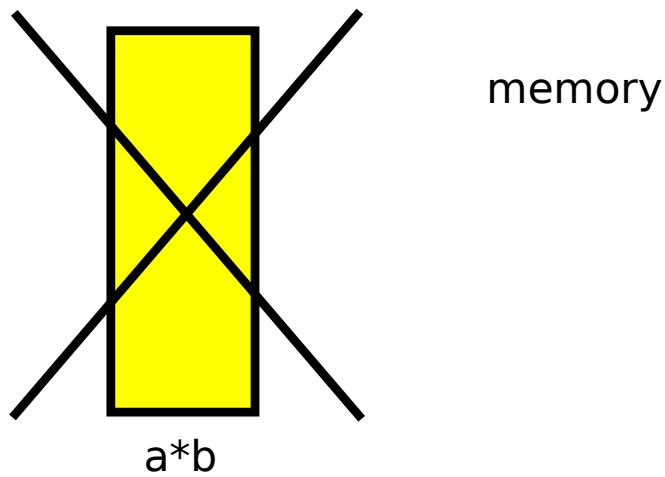
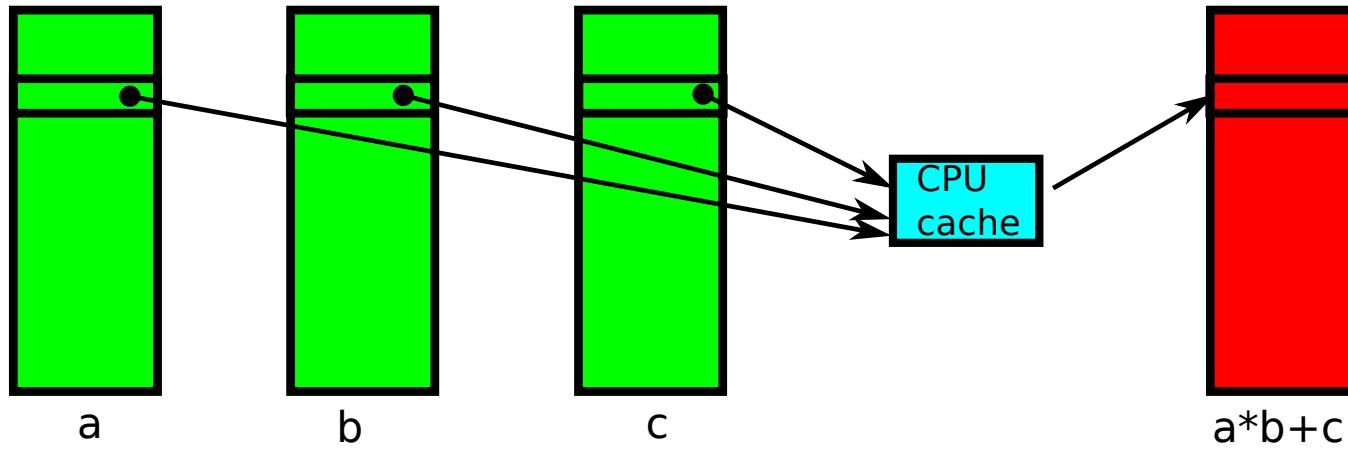
Time to evaluate polynomial (1 thread)



Computing " $a*b+c$ " with NumPy. Temporaries goes to memory.



Computing " $a*b+c$ " with Numexpr. Temporaries in memory are avoided.



Second Message of the Day

- Before spending too much time optimizing by yourself make you a favor:

Use the existing, powerful libraries out there

It is pretty difficult to beat performance professionals!

Optimal Containers for Big Data

The Need for a Good Data Container

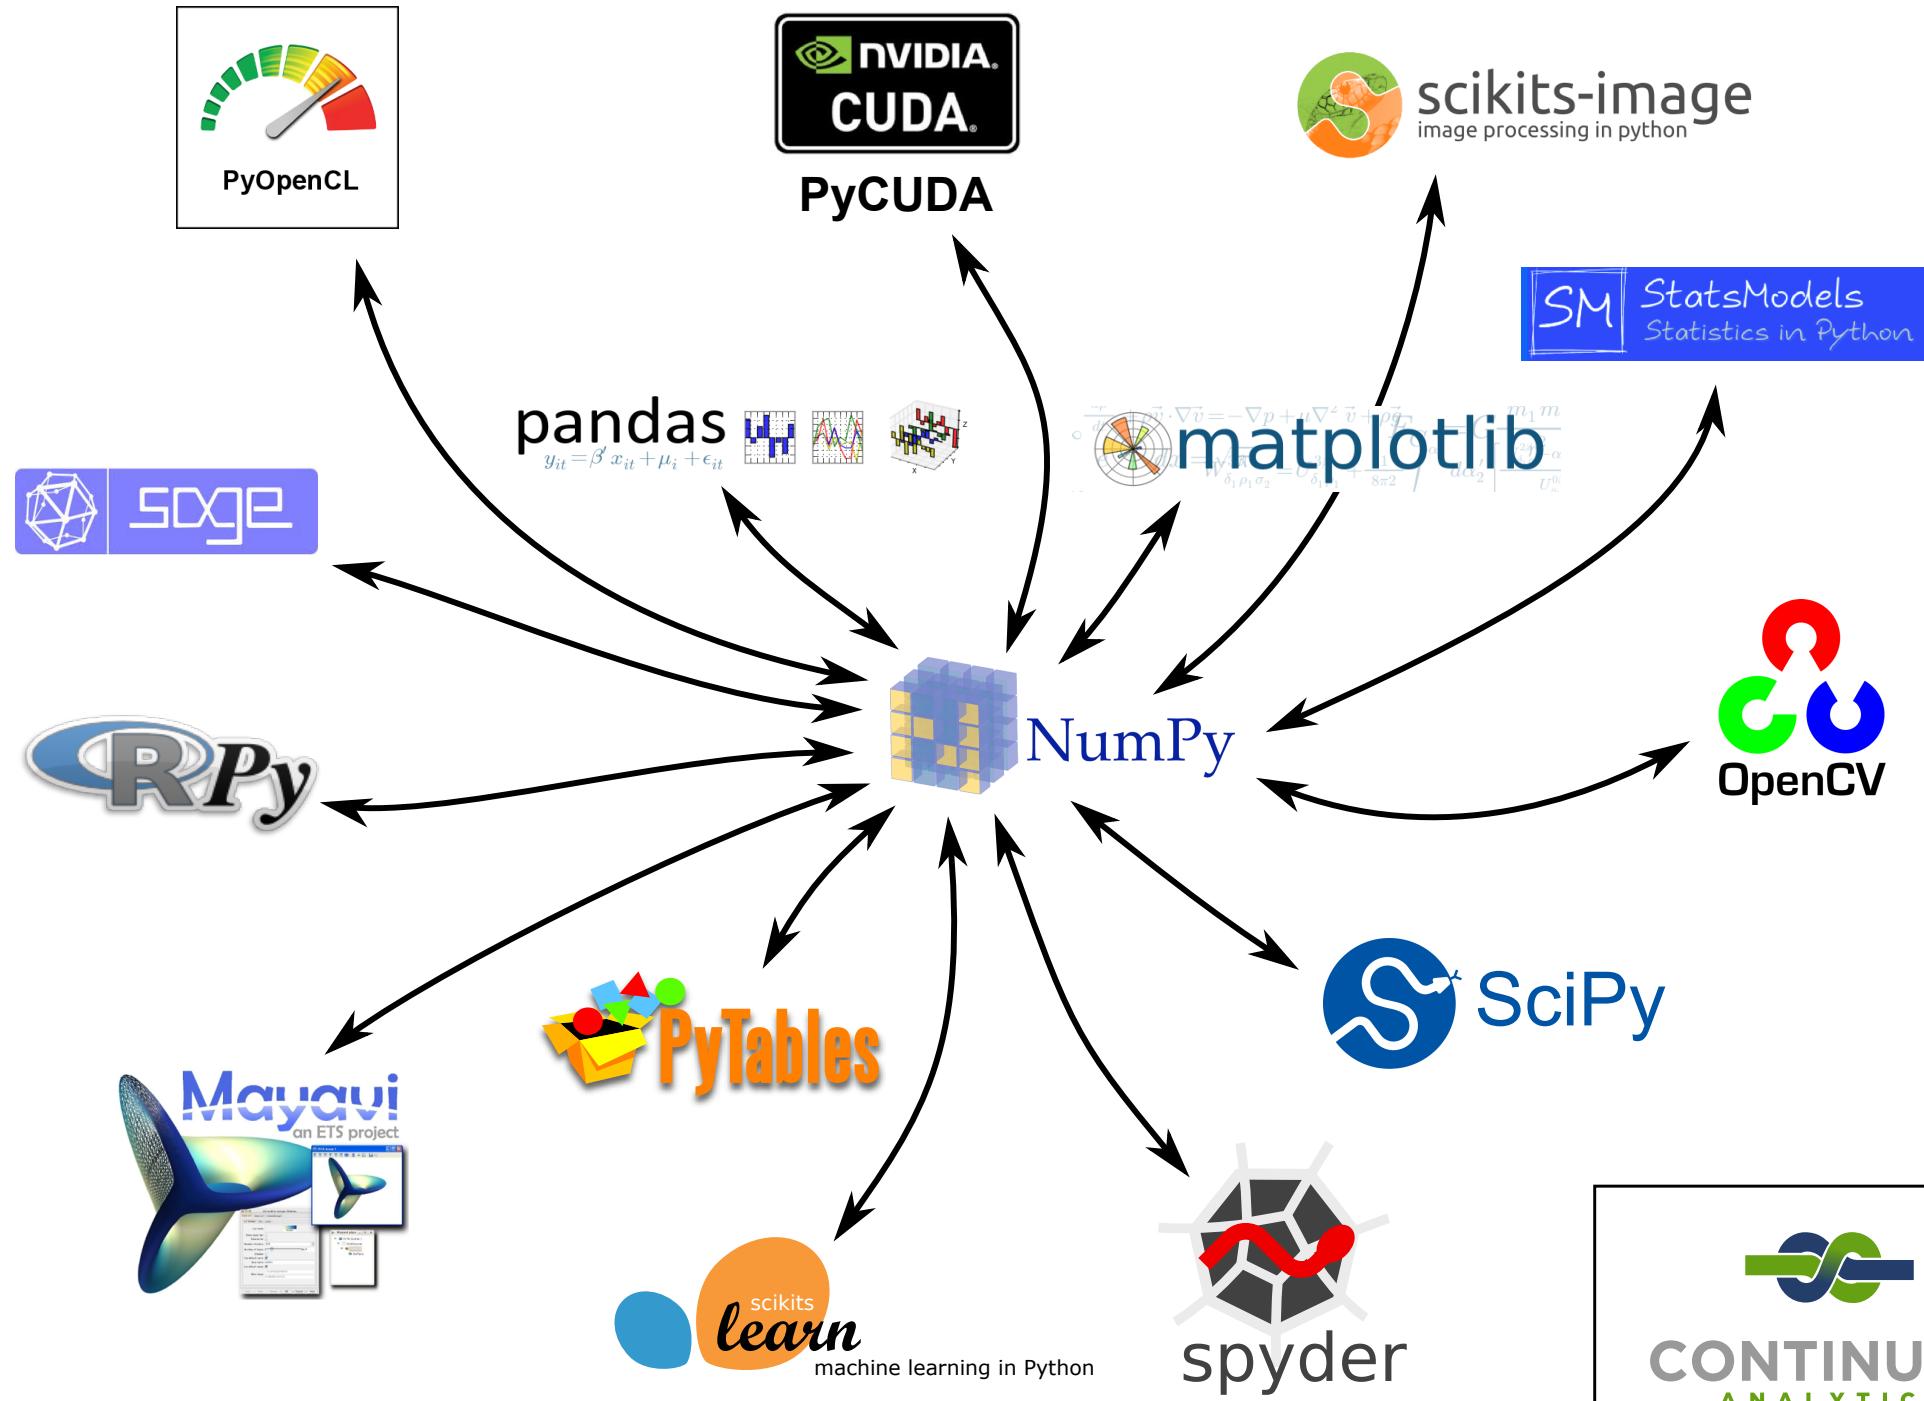
- Too many times we are focused on computing as fast as possible
- But we have seen how important data access is
- Hence, having an optimal data structure is critical for getting good performance when processing very large datasets

No Silver Bullet

- Unfortunately, there is not (and probably will never be) a fit-all data container
- We need to make our mind to the fact that we need to choose the ‘optimal’ container **for our case**

NumPy: A De Facto Data Container

NumPy is the standard de facto in-memory container for Big Data applications in the Python universe



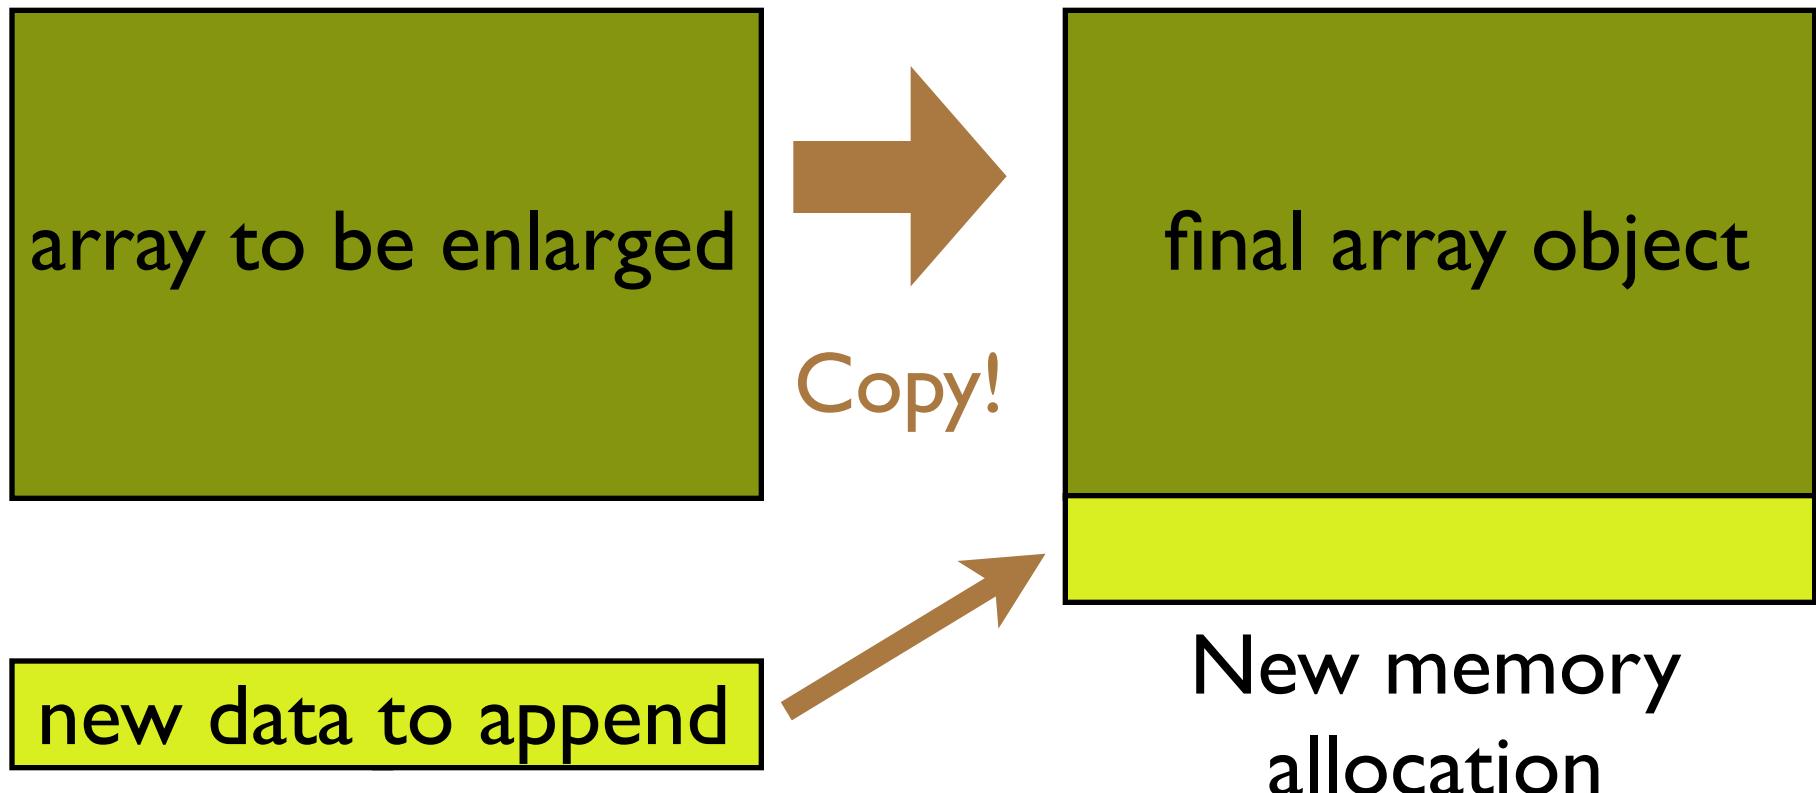
NumPy Advantages

- Multidimensional data container
- Efficient data access for many scenarios
- Powerful weaponry for data handling
- Efficient in-memory storage

Nothing Is Perfect

- The NumPy container is just great for many use cases
- However, it also has its own deficiencies:
 - Not efficient for appending data (so data containers tend to be static)
 - Cannot deal with compressed data transparently
 - Limited disk-based data support

Appending Data in Large NumPy Objects



- Normally a `realloc()` syscall will not succeed
- Both memory areas have to exist **simultaneously**

bcolz

Overcoming NumPy Limitations

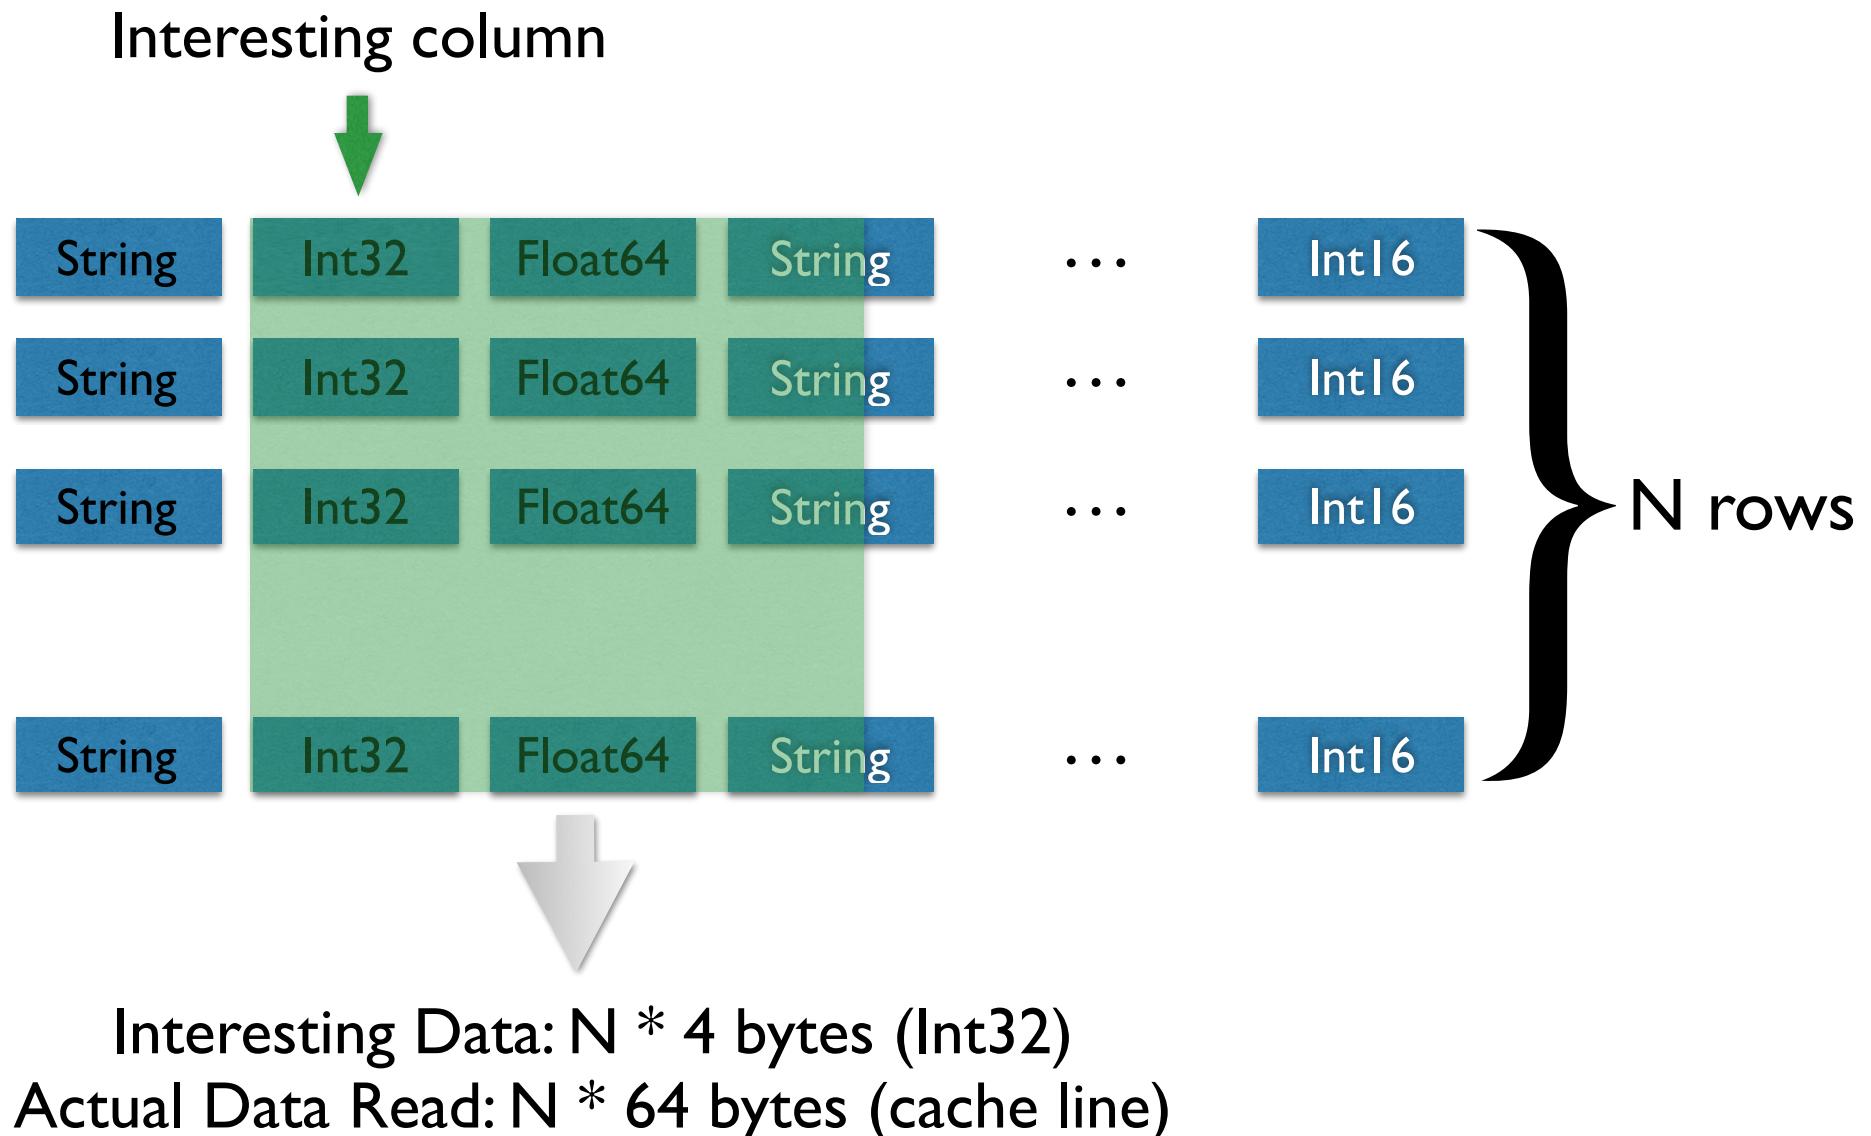
bcolz

- **Columnar, chunked, compressed data containers for Python**
- Offers **carray** and **ctable** container flavors
- Implements just a **few simple, but fast iterators** over the containers, supporting query semantics
- Uses the powerful **Blosc** compressor under the hood

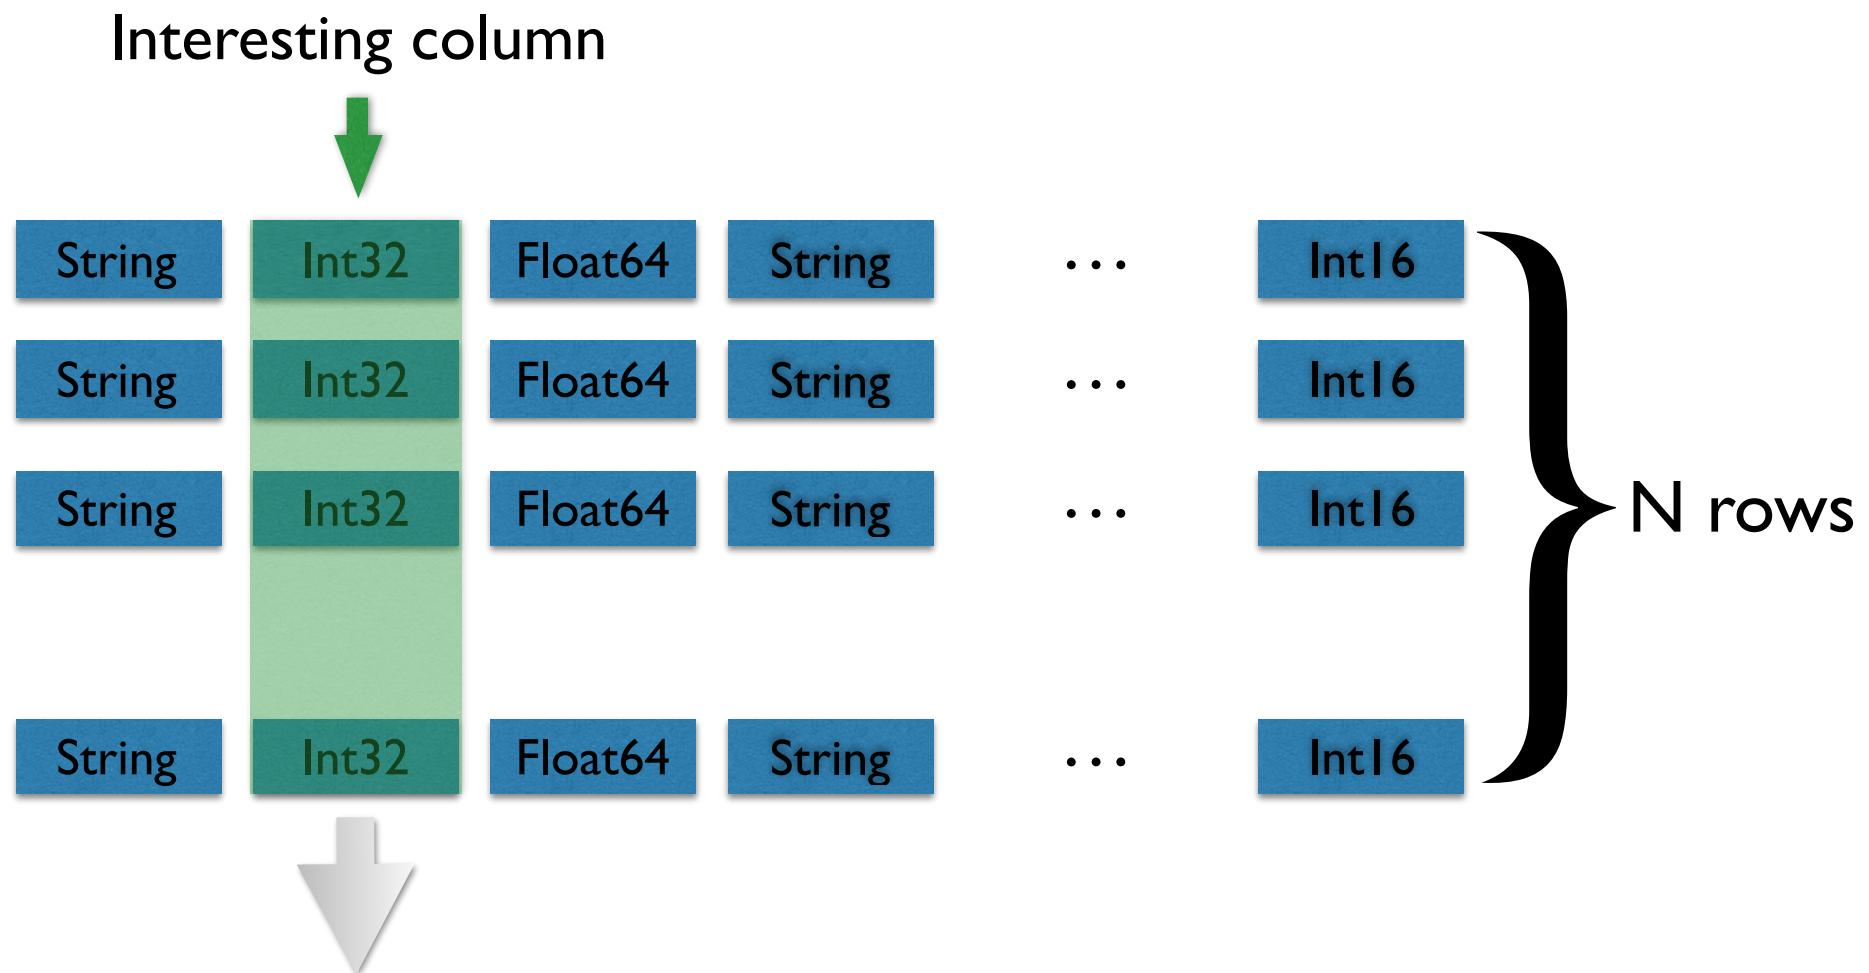
Why Columnar?

- When querying tabular data, only the interesting data is accessed
- Less I/O required

In-Memory Row-Wise Table



In-Memory Column-Wise Table



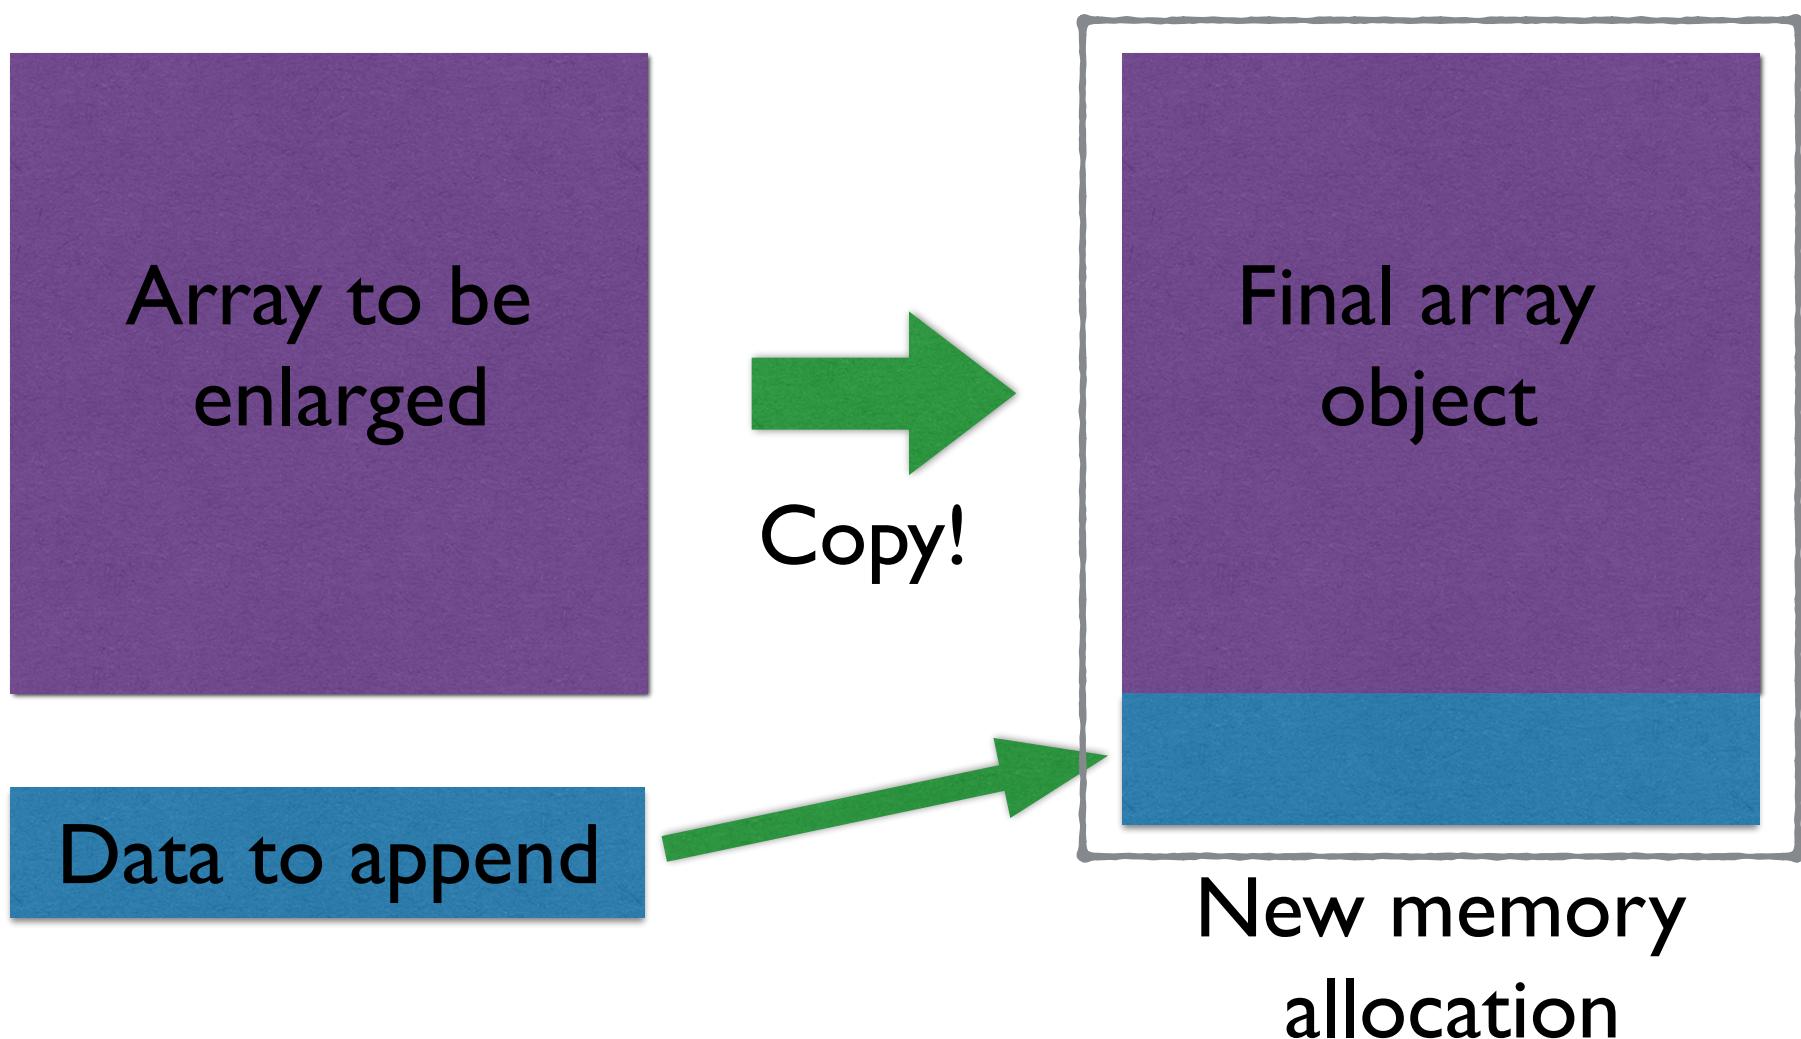
Interesting Data: $N * 4$ bytes (Int32)

Actual Data Read: $N * 4$ bytes (Int32)

Why Chunking?

- Chunking means more difficulty handling data, so why bother?
 - Efficient enlarging and shrinking
 - On-flight compression possible

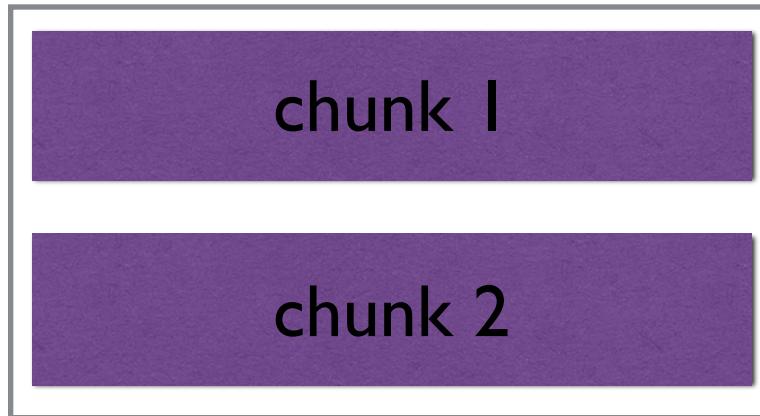
Appending Data in NumPy



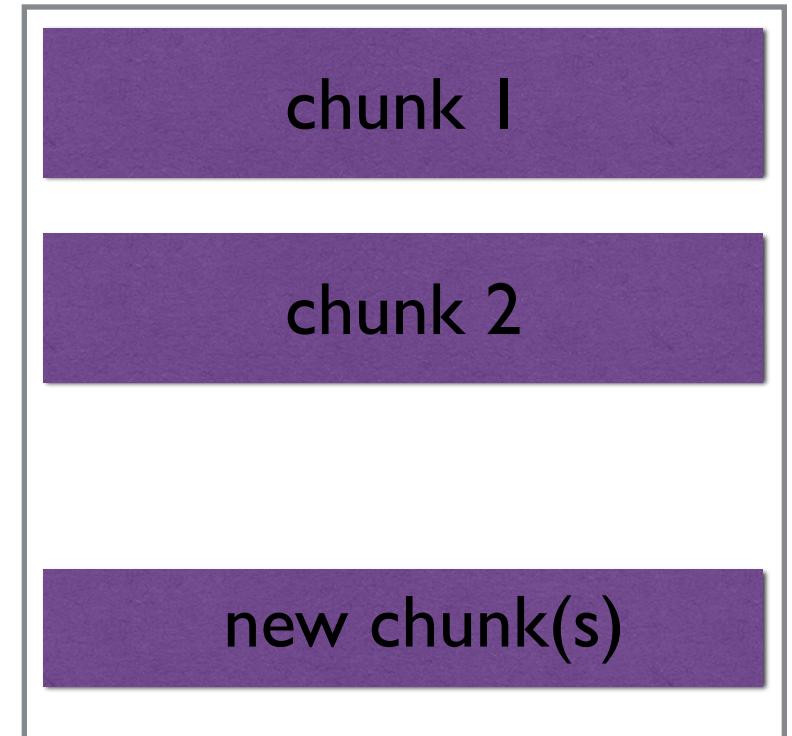
- Both memory areas have to exist **simultaneously**

Appending Data in bcolz

array to be enlarged



Final array object



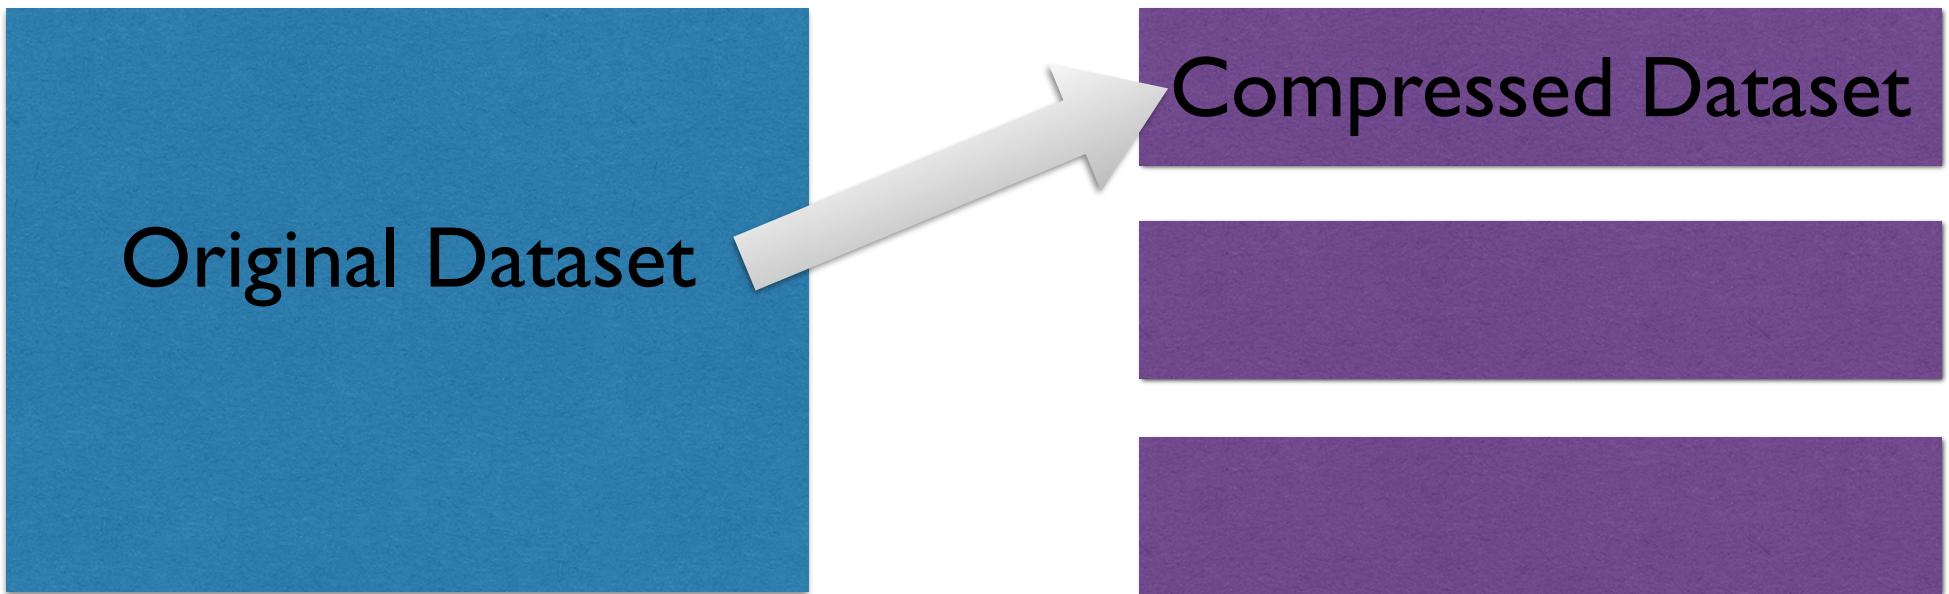
data to append

Compress

- Only a compression operation on new data is required

Why Compression (I)?

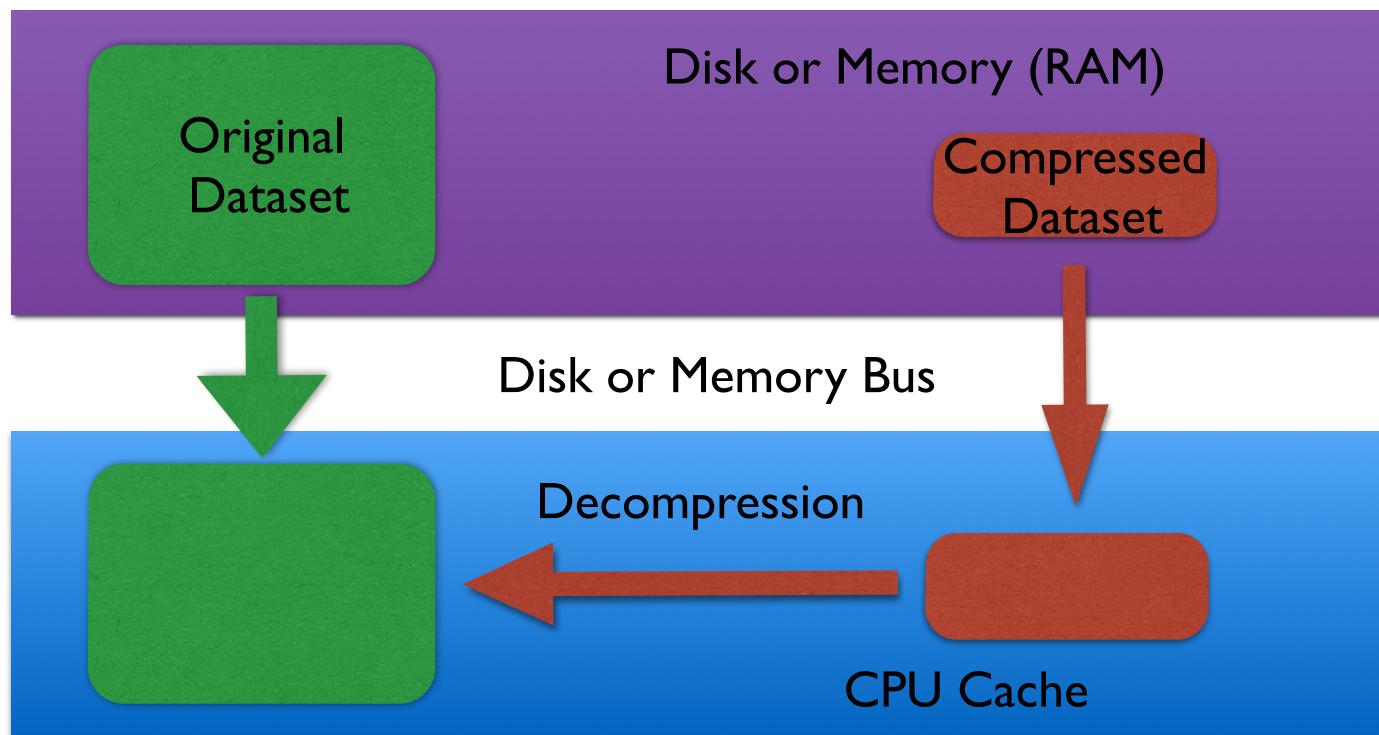
More data can be stored in the same amount of media



3x more data

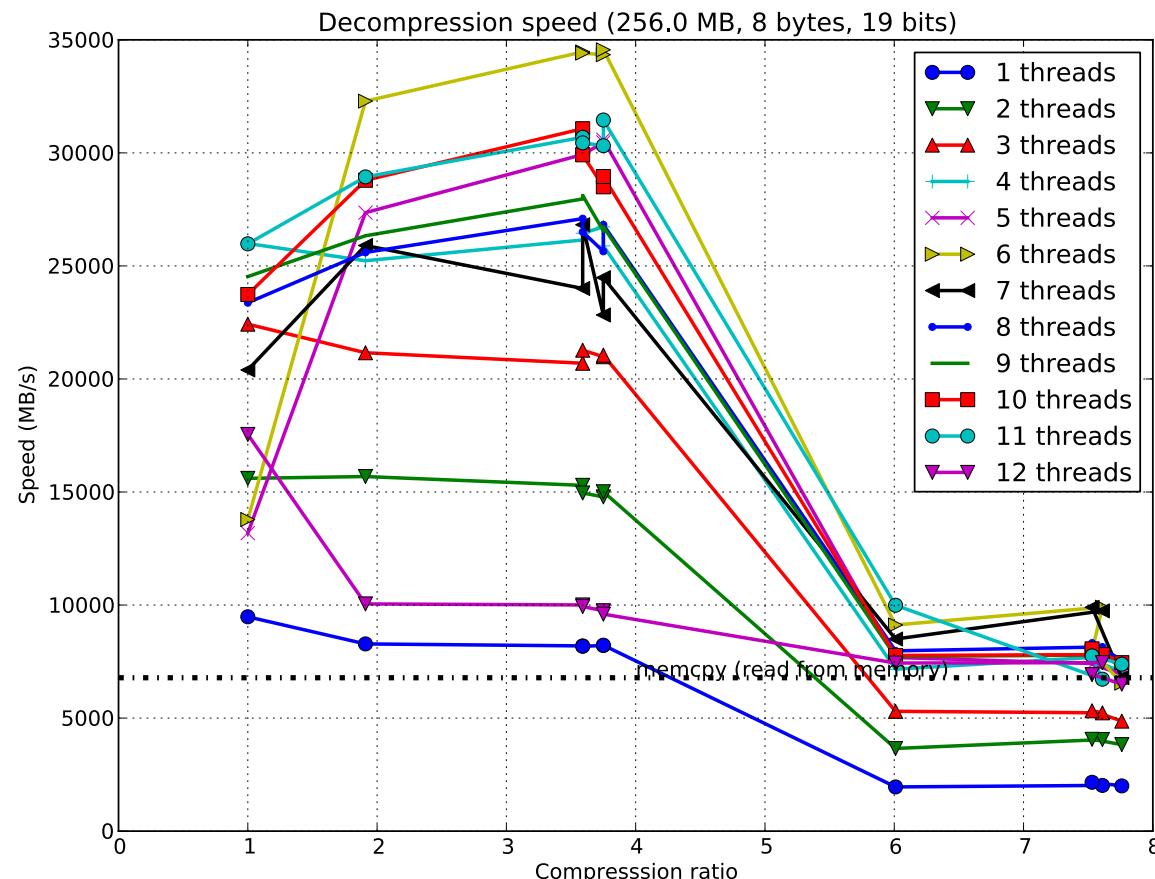
Why Compression (II)?

Less data needs to be transmitted to the CPU

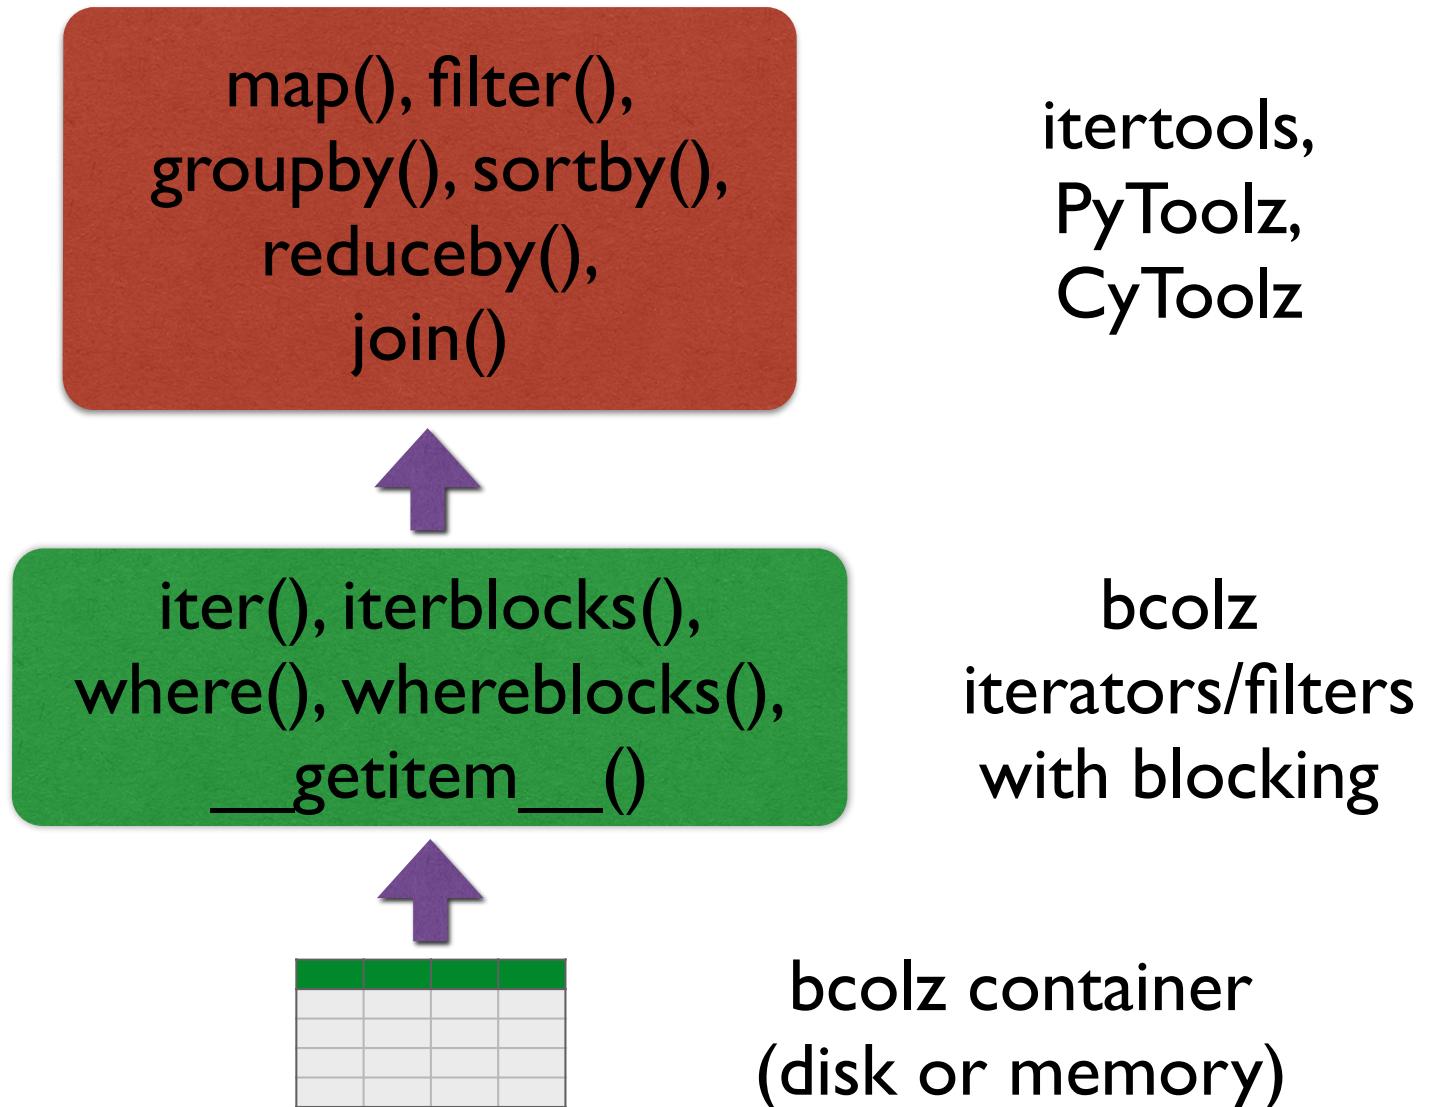


Transmission + decompression faster than direct transfer?

Blosc: Compressing Faster Than Memory Speed



Streaming Analytics with bcolz



Some Benchmarks With Real Data: The MovieLens Dataset

Materials in:

<https://github.com/Blosc/movielens-bench>

The MovieLens Dataset

- Datasets for movie ratings
- Different sizes: 100K, 1M, 10M ratings (the 10M will be used in benchmarks ahead)
- The datasets were collected over various periods of time

Querying the MovieLens Dataset

```
import pandas as pd
import bcolz

# Parse and load CSV files using pandas
...

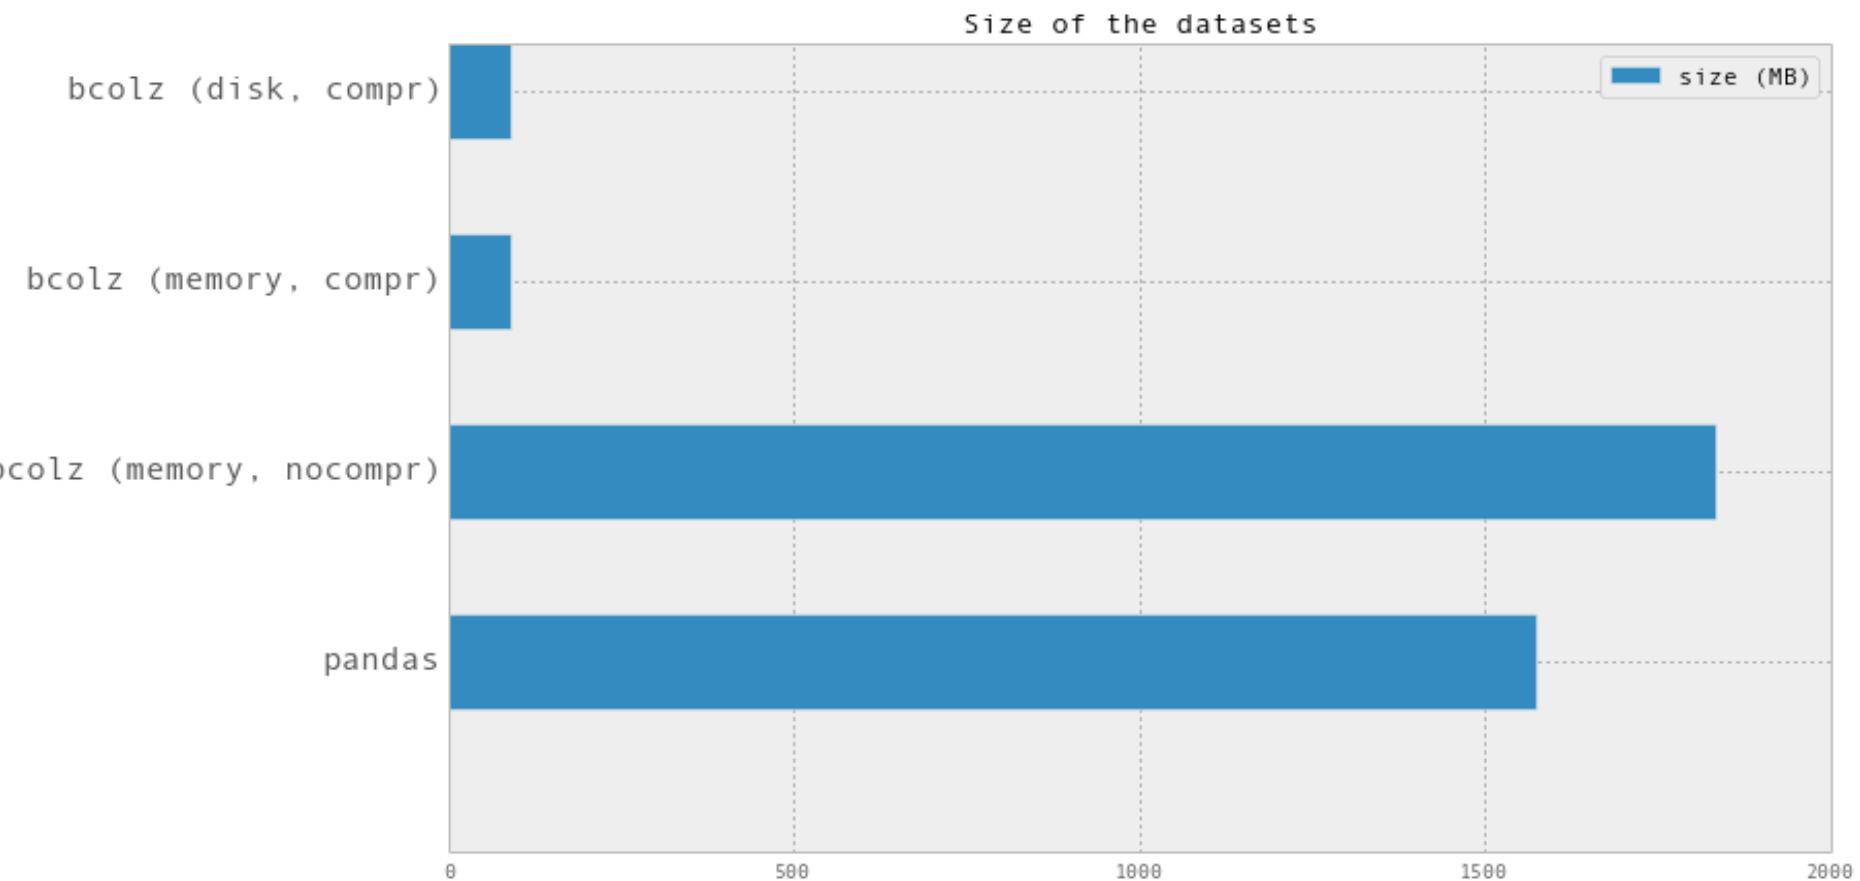
# Merge some files in a single dataframe
lens = pd.merge(movies, ratings)

# The pandas way of querying
result = lens.query("(title == 'Tom and Huck (1995)') & (rating == 5)")
['user_id']

# Get a ctable container
zlens = bcolz.ctable.fromdataframe(lens)

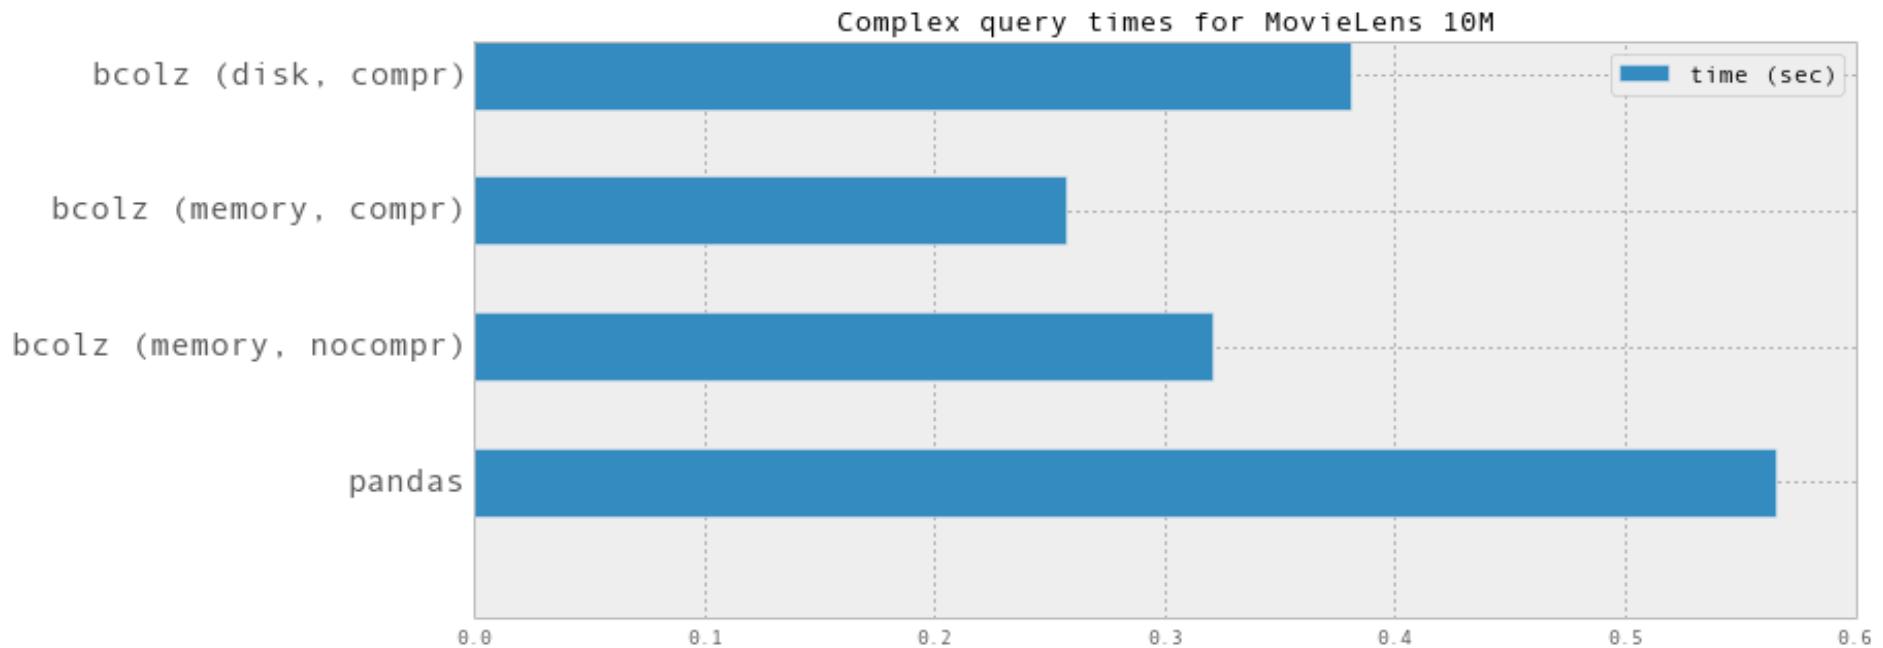
# The bcolz way of querying (notice the use of the `where` iterator)
result = [r.user_id for r in dblens.where(
    "(title == 'Tom and Huck (1995)') & (rating == 5)", outcols=['user_id'])]
```

Size of Datasets



- Compression means ~20x less space
- The uncompressed ctable is larger than pandas

bcolz Query Performance



- Compression leads to better query speeds (15% faster)
- Querying a disk-based ctable is fast!

Final Take Away Message for Today

- Big data is tricky to manage:

**Look for the optimal containers for your
data**

Spending some time choosing your appropriate data container can be a big time saver in the long run

Summary

- Nowadays you should be aware of the memory hierarchy for getting good performance
- Leverage existing memory-efficient libraries for performing computations optimally
- Use the appropriate data containers for your different use cases

References

- Why Modern CPUs Are Starving And What You Can Do About It
- bcolz: <http://github.com/Blosc/bcolz>
- Blosc ecosystem: <http://blosc.io>

Thank you!

Questions?

francesc@blosc.io

[@FrancescAlt](https://twitter.com/FrancescAlt)